



UNIVERSITY OF WESTERN AUSTRALIA

MASTER OF PHYSICS (COMPUTATIONAL PHYSICS)

**Quantum Circuits for Solving Differential Equations via
Physics-Informed Effective Hamiltonians and GQSP
Quantum Imaginary Time Evolution**

Dhyan Patel

School of Physics, Mathematics and Computing

Supervised by

Jingbo WANG

6 October 2026

Declaration

This is to certify that:

1. This thesis comprises of my original work.
2. Due acknowledgment has been made in the text to all other materials used.
3. This thesis consists of less than 60 pages inclusive of tables and figures but exclusive of references and appendices.

I authorise the Head of the Department of Physics to make a copy of this thesis to any person judged to have an acceptable reason for access to the information, i.e, for research, study or instruction.



Dhyan Patel

Student ID: 23340229

Date: 6 October 2026



Jingbo Wang

The University of Western Australia

Date: 6 October 2026

Acknowledgments

Thank you to my supervisor, Jingbo Wang, for her guidance and support throughout my research, particularly for her continual encouragement and belief in me.

I would also like to thank my family and friends for their support and encouragement throughout my studies.

Summary of Student Achievement

Abstract

Brief summary.....

Contents

1	Introduction	1
1.1	Differential Equations as a Target for Quantum Computation	1
1.2	Physics-Informed Hamiltonians and Where Their Cost Lies	1
1.3	Contributions and Scope	1
1.4	Outline	1
2	Background and Literature Review	2
2.1	Quantum Computation Preliminaries	2
2.1.1	Quantum Gates and Circuits	3
2.1.2	Complexity of Quantum Algorithms	4
2.1.3	Noisy Intermediate-Scale Quantum (NISQ) Era	4
2.2	Quantum Algorithms for Differential Equations	5
2.2.1	Linear-Systems and Linear-Combination Routes	5
2.2.2	Nonlinear Equations: Carleman and Related Linearisations	5
2.2.3	Variational Routes	5
2.2.4	Ground-State Routes and the Physics-Informed Lineage	5
2.2.5	Critical Assessment	5
2.3	Spectral Methods and Their Conditioning	5
2.3.1	Chebyshev Series, Galerkin Truncation and Differentiation	5
2.3.2	Banded Differentiation: the Ultraspherical Method	5
2.3.3	First-Order Reformulation and Least Squares	5
2.3.4	Fourier–Galerkin Discretisation of Periodic Flows	5
2.4	Block Encoding, Linear Combinations of Unitaries and Qubitisation	5
2.5	Quantum Signal Processing and Its Generalisation	5
2.6	Ground-State Preparation by Filtering	5
2.7	Readout and Amplitude Amplification	5
2.8	Synthesis	5
3	The Physics-Informed Effective Hamiltonian Method	6
3.1	Overview	6
3.2	The Chebyshev Latent Space Model	6
3.3	The Residual, the Hamiltonian and Its Constraints	7
3.4	The Published Circuits, Filter and Readout	10
3.5	The Doubled-Space Extension to Nonlinearity	12
3.6	Critical Assessment	13
4	The Standard Form: Encoder and Hamiltonian Costs	15
4.1	The Published Encoding, Rebuilt	15
4.2	An Exact Factorisation of the Differentiation Matrix	16
4.3	Circuits and Cost	17

4.4	What Remains: The Standard-Form Hamiltonian	19
5	The Descriptor Reformulation	21
5.1	The Idea	21
5.2	The Banded Factorisation	22
5.3	Two Layouts and Kernel Equivalence	23
5.4	Constraints, Partial Differential Equations and Coupled Fields	25
5.5	The Gap of a Carried Derivative	26
5.6	Block Rescaling	27
5.7	The Block Encoding: One Flat Linear Combination	28
5.8	Cost of the Descriptor Encoding	30
5.9	Consequences for the Spectrum and the Filter	31
5.10	Summary	32
6	Polynomial Nonlinearity	33
6.1	Two Treatments	33
6.2	The Doubled Space: Exact Encoding, Degenerate Kernel	33
6.3	The Carleman Lift in Space–Time	33
6.4	Products and Couplings on a Circuit	33
6.5	Fourier–Galerkin Fluids: Burgers and Navier–Stokes	33
6.6	Preparation Versus Projection	33
7	Preparing and Verifying the Ground State	34
7.1	Overview and Model Configuration	34
7.2	Problems and Model Inputs	34
7.3	The Classical Reference	34
7.4	From Stacked Residual to Quantum Walk	34
7.5	The Filter and Its Parameters	36
7.6	Verification and the Error Budget	37
7.7	Reading Out the Solution	37
7.8	Design of the Numerical Experiment	37
7.9	Software and Computation	37
8	Numerical Study	38
8.1	Reproducing Wu et al.	38
8.2	The Pipeline Is Faithful	40
8.3	Where the Cost Lies	40
8.4	The Descriptor’s Gap and Its Rescaling	40
8.5	Nonlinear Problems	40
8.6	Case Study: Two-Dimensional Navier–Stokes	40
8.7	Summary of Findings	40

9	Discussion	43
9.1	Why the Descriptor Form Wins, and What It Costs	43
9.2	End-to-End Complexity	43
9.3	What the Reproduction Says About the Published Method	43
9.4	Where the Construction Applies	43
9.5	Feasibility on Fault-Tolerant Hardware	43
9.6	Limitations	43
10	Conclusions and Future Work	44
10.1	Conclusions	44
10.2	Future Work	44
11	References	45
	Appendices	50

List of Figures

1	The odd-offset shifts are one subtraction. PREPARE (left) puts the offset register d in a uniform superposition over the $N/2$ odd offsets, by an X on its lowest qubit and Hadamards on the rest. SELECT is a ripple-carry subtractor [10] that replaces the system register p by $p - d$ and sets the borrow flag when $p < d$. PREPARE [†] (right) with post-selection of d and the borrow flag on $ 0\rangle$ keeps $\frac{2}{N} k\rangle\langle p $ for every odd $d = p - k \geq 1$, and the wrapped terms vanish exactly. The block is $U_{\text{odd}}/(N/2)$, so $\alpha = N/2$, on n system and $n + 2$ ancilla qubits.	17
2	Dense against banded: the two residuals of the same equation.	22
3	The descriptor block encoding: one PREPARE/SELECT/PREPARE [†] over the whole residual.	29
4	Model configuration: from problem specification to decoded field.	34
5	The GQSP filter circuit.	36
6	Printed choices distort the paper's solutions at the printed resolution.	38
7	Every encoding is verified to its mode's tolerance at every size tested.	39
8	The prepared state reproduces the analytic solution to within the predicted error budget.	39
9	The encoder wall and the Hamiltonian wall are separate gaps.	40
10	Three representative problems: where the comparison, quoted as built, reverses at the ideal encoding bound.	41
11	Block rescaling restores an $O(1)$ gap on one axis, and stops at a double zero or on two axes.	41
12	Three representative problems, solved in both forms.	42
13	Carleman truncation converges geometrically once time is resolved; the doubled space only projects; on Burgers the error tracks a/ν , not ρ	42
14	Navier–Stokes is reached by exact emulation and a resource estimate, not by circuit simulation.	43
15	The layers of the <code>pihm</code> package and the one direction in which they may import.	67

List of Tables

1	Quantum approaches to differential equations: what each assumes, what it costs, and where it fails.	5
2	The paper’s claims for its block encodings, tested against the rebuilt circuits. Figs 2 and 8 hold as printed. In Fig. 9, the gate count and the error hold, the qubit count is two higher, and the prefactor is valid but loose against $\ \mathbb{G}\ _2$ by a factor that grows as N^3 (Table 3); the polynomial gate count does not make the encoding efficient. Classical exact-mode build; circuit counts as built, error to leading order.	16
3	Subnormalisation of the differentiation-matrix encodings, relative to $\ \mathbb{G}\ _2$. The structured encoding of Proposition 2 stays within a factor 2.12 of the norm, and the published prefactor of Equation (3.11) grows as N^3 . The full table to $n = 10$ is Table 14. Classical exact-mode computation.	18
4	Where the descriptor’s degree changes. Norms follow from the block weights, gaps are measured. The block weight ζ is the a priori rule of Section 5.6.	31
5	Which nonlinear route prepares the solution, and which only projects onto a kernel that contains it.	34
6	The shared constraint layer: every constraint row, in either form, is one of four kinds.	35
7	The benchmark ladder. Each problem adds one capability to those below it; R2–R6 are the paper’s Figs 3–7.	35
8	Operator-verification modes: what each compares, up to what size, and to what tolerance.	37
9	Acceptance criteria for every state-preparation run.	37
10	The paper’s printed scale factors against the exact solution’s $\ b\ ^2$ (Equation (8.1)) and the paper-faithful ground state at the printed n . [†] : the ground state is not trusted (a degenerate kernel, or a gap ratio below 100). Classical computation.	38
11	Where the cost lies: the paper’s regime under its published encoding and under the structured exact encoding, against the descriptor form, for Fig. 3b ($\{1, 4, 4\}$, $f(-1) = 0$). Tier: T2 where measured, T3 beyond.	40
12	The gates used in this thesis.	53
13	The published encoding of $\mathbb{G}$ against the paper’s claims. Classical exact-mode measurement.	58
14	Subnormalisation of the structured and published encodings of $\mathbb{G}$, relative to $\ \mathbb{G}\ _2$	59
15	Accuracy of the classical reference, $1 - \cos $ against the exact coefficients, for the standard form by $\text{eigh}(R^\top R)$ and by SVD of R , and for the descriptor form.	62
16	The linear benchmark problems, R1–R5, with the paper’s printed invariant constraints.	63

17	The nonlinear and coupled benchmark problems, R6–R10.	63
18	State preparation by problem and form. Tier as marked.	64
19	The paper’s solutions at its printed resolution: L^∞ error of the paper-faithful and corrected variants, and the corrected variant’s best over n . Classical computation.	65
20	Which printed choice causes the Fig. 5 error: the paper-faithful variant over $n = 4, \dots, 10$, each correction alone, and both together.	65
21	The claims of Wu et al. tested in this thesis, and their verdicts.	66
22	Dropping the top k rows of the equation block (the Chebyshev tau method) gives an exact kernel with accuracy comparable to the paper’s least squares. . .	67
23	Resource classes for the Setonix campaigns.	68
24	The test tiers of the <code>pihm</code> package.	68
25	Deviations from the research proposal of May 2025.	79

Notation and Conventions

A quick reference for symbols used across the thesis; each is properly introduced at its first use in the text below.

Discretisation

Symbol	Meaning	Introduced
n	Qubits per Chebyshev axis (the resolution).	Section 3.2
$N = 2^n$	Truncation dimension of one Chebyshev axis.	Section 3.2
k	Order of an ordinary differential equation.	Section 3.3
p	Degree of a polynomial nonlinearity.	Section 6.3
n_f	Qubits in the field (block) register of the descriptor form.	Section 5.3
$\dim w$	Padded descriptor state dimension, $2^{\lceil \log_2(k+1) \rceil} N$.	Appendix G.2

The Chebyshev model

Symbol	Meaning	Introduced
T_ℓ	Chebyshev polynomial of the first kind, degree ℓ .	Section 3.2
$ \tau(x)\rangle$	The raw Chebyshev feature state.	Section 3.2
ψ, η, η_e	Normalised coefficient state; its scale; $\eta_e = \ b\ ^2$.	Section 3.2
$\hat{f}^{(j)}$	Coefficients of $f^{(j)}$: one block of the descriptor state.	Section 5.3
w	The descriptor state $(\hat{f}^{(0)}; \dots; \hat{f}^{(k)})$.	Section 5.3

Operators and constraints

Symbol	Meaning	Introduced
$\mathbb{G}$	Differentiation matrix, $f \mapsto f'$, in coefficient space: dense, strictly upper triangular, $\ \mathbb{G}\ _2 = \Theta(N^2)$ (Wu et al. print it as $\mathbb{G}_n^\top$).	Section 3.3
$\tilde{\mathbb{B}}, \tilde{\mathbb{D}}$	Banded factors of the differentiation matrix, $\mathbb{G} = \tilde{\mathbb{B}}^{-1}\tilde{\mathbb{D}}$.	Section 5.2
$\mathbb{M}_{x^p}, \mathbb{M}_a$	Multiplication by x^p (the exact lift, $2N \times N$) and by a smooth $a(x)$.	Section 3.3
$\mathfrak{m}_1, \mathfrak{m}_x$	Product-to-sum folds.	Section 3.5
$\mathbb{B}(x), \mathbb{D}^{(0)}(x_s)$	Rank-one zero (invariant) and collapse (regular) constraint rows.	Section 3.3
x_z, x_m, x_s	Constraint points: a zero, an extremum, a source datum.	Section 3.3
$R_0, U_{\text{odd}}, \Lambda$	Factors of the structured exact encoding of $\mathbb{G}$.	Section 4.2
$\mathbb{U}$	The Chebyshev–Gauss transform used by the nodal fold.	Section 6.4
$\ \mathbb{G}\ _S$	The paper’s norm $\max(\ \mathbb{G}\mathbb{G}^\top\ _1, \ \mathbb{G}^\top\mathbb{G}\ _1)$.	Section 3.4

Residuals and Hamiltonians

Symbol	Meaning	Introduced
$A_{\text{std}}, A_{\text{desc}}$	Residual of the standard form; of the descriptor form.	Section 3.3
$H_{\text{std}}, H_{\text{desc}}$	Effective Hamiltonian of each form; bare H only in regime-independent statements.	Section 3.3
$R = [R_1; \dots; R_L]$	Stacked residual, with $H = R^\dagger R$.	Section 7.4
$H_\otimes$	The doubled-space Hamiltonian.	Section 6.2

Block encoding, filtering and state preparation

Symbol	Meaning	Introduced
$\alpha, \alpha_R, \alpha_H = \alpha_R^2$	Subnormalisation: generic, of R , of H .	Section 2.4
V, W	Stacked-residual reflection; qubitised walk.	Section 7.4
$H' = 2H/\alpha_R^2 - I$	The normalised Hamiltonian that V block-encodes.	Section 7.4
Δ	Spectral gap $\lambda_1 - \lambda_0$.	Section 7.5
β	Imaginary time.	Section 7.5
d, q_k	GQSP degree; filter coefficients.	Section 7.5
ε	Target infidelity (10^{-6}).	Section 7.5
φ_0, γ	Initial state; its overlap $ \langle \psi_{\text{ref}} \varphi_0 \rangle $ with the reference.	Section 7.5
p_G	GQSP success probability.	Section 7.5
ζ	Block-rescaling constant.	Section 5.6
$\epsilon_{\mathbb{G}}, \mu$	Entry error and per-bit coefficient of the published encoding of $\mathbb{G}$.	Appendix D.2

Nonlinearity

Symbol	Meaning	Introduced
K	Carleman truncation order.	Section 6.3
M	Dimension of the base state u .	Section 6.3
ρ	Nonlinearity-to-dissipation ratio; the Carleman precondition is $\rho < 1$.	Section 6.3
$F_1, F_2, \mathbb{A}_K$	Linear and quadratic terms; the lifted Carleman generator.	Section 6.3

Acronyms

LCU linear combination of unitaries.

ODE ordinary differential equation.

PDE partial differential equation.

QITE quantum imaginary-time evolution.

QSVT quantum singular value transformation.

1 Introduction

1.1 Differential Equations as a Target for Quantum Computation

1.2 Physics-Informed Hamiltonians and Where Their Cost Lies

1.3 Contributions and Scope

1.4 Outline

2 Background and Literature Review

2.1 Quantum Computation Preliminaries

In classical computing, the fundamental unit of information is the bit, which can take one of two values: 0 or 1.

The fundamental unit of information in quantum computing is the qubit, which can exist in a superposition of states, allowing it to encode both 0 and 1 simultaneously. This property enables quantum computers to perform certain computations more efficiently than classical computers.

Definition 1 (State Space): The set of all quantum states that a quantum system can occupy is a complex vector space with inner product structure known as a Hilbert space (denoted $\mathcal{H}$). A finite dimensional Hilbert space is isomorphic to $\mathbb{C}^N$ ($\mathcal{H} \cong \mathbb{C}^N$) for some $N \in \mathbb{N}$, where N is the dimension of the Hilbert space. A quantum state is a unit vector in this Hilbert space, represented in dirac notation as $|\psi\rangle \in \mathcal{H}$, with $\langle\psi|\psi\rangle = 1$ (normalisation condition). A quantum state $|\psi\rangle$ can be expressed as a linear combination of basis states $\{|i\rangle\}_{i=0}^{N-1}$, where $|i\rangle$ is the i^{th} basis state of the Hilbert space.

$$|\psi\rangle = \sum_{i=0}^{N-1} \psi_i |i\rangle = \begin{pmatrix} \psi_0 \\ \psi_1 \\ \vdots \\ \psi_{N-1} \end{pmatrix} \quad (2.1)$$

Its hermitian conjugate (or adjoint) is denoted as $\langle\psi| = |\psi\rangle^\dagger$. Note that for some $e^{i\theta}$ for $\theta \in [0, 2\pi)$, the states $|\psi\rangle$ and $e^{i\theta}|\psi\rangle$ are physically indistinguishable, and thus correspond to the same physical quantum state. This is known as a global phase factor.

Definition 2 (The Qubit): A quantum bit (qubit) is a two-state quantum system, characterised by $|\psi\rangle \in \mathbb{C}^2$.

$$|\psi\rangle = \alpha|0\rangle + \beta|1\rangle = \begin{pmatrix} \alpha \\ \beta \end{pmatrix} \quad (2.2)$$

where $\alpha, \beta \in \mathbb{C}$, $|0\rangle$ and $|1\rangle$ are the computational basis states, and $|\alpha|^2, |\beta|^2$ are the probabilities of measuring the qubit in the $|0\rangle$ and $|1\rangle$ states, respectively. The normalisation condition requires that $|\alpha|^2 + |\beta|^2 = 1$.

Definition 3 (Composite Systems): The state space of a composite quantum system is the tensor product of the state spaces of its subsystems. Consider two quantum states $|\psi\rangle_A \in \mathcal{H}_A$

and $|\psi\rangle_B \in \mathcal{H}_B$, the composite state $|\psi\rangle_{AB}$ is given by:

$$|\psi\rangle_{AB} = |\psi\rangle_A \otimes |\psi\rangle_B \in \mathcal{H}_A \otimes \mathcal{H}_B \quad (2.3)$$

where $\otimes$ denotes the tensor product and is often hidden by convention. The dimension of the composite state space is the product of the dimensions of the individual state spaces. If a composite system cannot be written as a tensor product of its subsystems, it is said to be entangled. Entanglement is a unique feature of quantum mechanics that has no classical counterpart and plays a crucial role in quantum information processing.

We often consider the state of a composite system of n qubits, which is represented in a 2^n -dimensional Hilbert space $\mathcal{H} \cong \mathbb{C}^{2^n}$.

2.1.1 Quantum Gates and Circuits

Definition 4 (Quantum Algorithm): A quantum algorithm is a sequence of operations that manipulates qubits to perform a specific computation.

The goal of any quantum algorithm is to manipulate quantum states in a controlled manner to achieve a desired outcome. This is accomplished through the application of quantum gates, which are unitary operations that transform quantum states. Quantum gates can be combined to form quantum circuits, which represent the sequence of operations performed on qubits. Each quantum gate can be implemented using physical operations on hardware, and the design of efficient quantum algorithms is a central challenge in quantum computing research.

Definition 5 (Unitary Evolution): The evolution of a closed quantum system between two states is described by unitary transformations $U \in \mathbb{C}^{2^n \times 2^n}$, where n is the number of qubits in the system.

$$|\psi\rangle_{final} = U |\psi\rangle_{initial}, \quad U^\dagger U = U U^\dagger = I \quad (2.4)$$

where $U^\dagger$ is the hermitian conjugate of U and I is the identity operator. Unitary transformations preserve the inner product and thus the norm of quantum states, ensuring that states have a physical interpretation as probability amplitudes.

These unitary transformations are known as quantum gates, and each gate serves a specific purpose in manipulating quantum states.

The quantum circuit diagram offers a succinct and intuitive representation of quantum algorithms, analogous to Feynman diagrams in quantum field theory.

Time flows left to right, and each wire (horizontal line) represents a qubit or a quantum register. Gates are placed sequentially on the wires indicating operations applied to qubits. An example is the Hadamard gate, which creates superposition, and the CNOT gate, which entangles qubits. A full list of quantum gates and their corresponding circuit symbols can be found in .

May
need
to re-
write

Create
a
table
in the
ap-
pendix

2.1.2 Complexity of Quantum Algorithms

Definition 6 (Gate Complexity): The total number of quantum gates used in a quantum circuit.

Definition 7 (Circuit Depth): The maximum number of gates applied sequentially on any qubit in a quantum circuit.

Definition 8 (Ancillary Qubits): Additional qubits introduced into a quantum circuit beyond the primary input qubits, typically initialised to a baseline state (such as $|0\rangle$) to assist with intermediate calculations, temporary storage, or entanglement generation.

Definition 9 (Big- $\mathcal{O}$ Notation): Big- $\mathcal{O}$ notation is a mathematical notation used to describe the asymptotic behaviour of functions.

$$\mathcal{O}(f(n)) = \{g(n) : \exists c > 0, n_0 \in \mathbb{N} \text{ such that } 0 \leq g(n) \leq cf(n) \text{ for all } n \geq n_0\} \quad (2.5)$$

In the context of quantum algorithms, it is used to express the scaling of gate complexity, circuit depth, and space requirements with respect to the number of qubits n .

For example, if a quantum algorithm has a gate complexity of $\mathcal{O}(n^2)$, it means that the number of gates required grows quadratically with the number of qubits.

A quantum circuit is said to be **efficient** if its circuit depth, gate complexity, and the number of ancillary qubits all scale polynomially with the number of qubits n .

2.1.3 Noisy Intermediate-Scale Quantum (NISQ) Era

Write this section. Use Josh's as a guide.

Table 1: Quantum approaches to differential equations: what each assumes, what it costs, and where it fails.

Approach	Assumptions	Query cost	Output	Failure mode
Linear systems (HHL family)				
Linear combination of Hamiltonian sim- ulations				
Carleman linearisa- tion				
Variational				
Physics-informed				
Hamiltonian [1]				
This work				

2.2 Quantum Algorithms for Differential Equations

2.2.1 Linear-Systems and Linear-Combination Routes

2.2.2 Nonlinear Equations: Carleman and Related Linearisations

2.2.3 Variational Routes

2.2.4 Ground-State Routes and the Physics-Informed Lineage

2.2.5 Critical Assessment

2.3 Spectral Methods and Their Conditioning

2.3.1 Chebyshev Series, Galerkin Truncation and Differentiation

2.3.2 Banded Differentiation: the Ultraspherical Method

2.3.3 First-Order Reformulation and Least Squares

2.3.4 Fourier–Galerkin Discretisation of Periodic Flows

2.4 Block Encoding, Linear Combinations of Unitaries and Qubit-isation

2.5 Quantum Signal Processing and Its Generalisation

2.6 Ground-State Preparation by Filtering

$$d = \Theta\left(\sqrt{\frac{\alpha_H}{\Delta}} \log \frac{1}{\varepsilon}\right), \quad \alpha_H = \alpha_R^2. \quad (2.6)$$

2.7 Readout and Amplitude Amplification

2.8 Synthesis

3 The Physics-Informed Effective Hamiltonian Method

3.1 Overview

This chapter reports the physics-informed effective Hamiltonian method of Wu, Paine, Philip, Gentile and Kyriienko [1] as published, in the register of reviewed literature, and then judges it. Nothing in Sections 3.2 to 3.4 and the first two paragraphs of Section 3.5 is a result of this thesis; the thesis’s own reproduction of the method is reported in Section 8.1. The judgement is ours and is confined to the last paragraph of Section 3.5 and to Section 3.6.

The method, in the form the paper gives it and which this thesis calls the standard form (Section 3.3), turns a differential equation into a ground-state problem in three steps. First, the unknown function is stored as a vector of Chebyshev coefficients, and that vector is treated as the amplitude vector of an n -qubit state $|\psi\rangle$, so that the value of the function at any point is an overlap between $|\psi\rangle$ and a known feature state. Second, the differential equation and every boundary or initial condition are written as matrices that annihilate $|\psi\rangle$ when it holds the coefficients of the solution. The equation $f'' + f = 0$, the paper’s introductory example, becomes $(\mathbb{G}^2 + I)|\psi\rangle = 0$, where $\mathbb{G}$ is the matrix that differentiates a Chebyshev series. Third, the squared violations of all these conditions are summed into one quadratic form $\langle\psi|H|\psi\rangle$, which defines a positive semi-definite Hamiltonian H whose ground state is the solution. Solving the equation is then the preparation of the ground state of H , which needs a block encoding of H and a ground-state filter, and the field is read back through the same feature map that defined the model. The lineage runs from differentiable quantum circuits [2], through the Chebyshev latent space [3], to the effective Hamiltonian.

The paper’s case for the method is efficiency. It reports that each term of the Hamiltonian block-encodes with a number of gates polynomial in n , and it expects the preparation time to scale inversely with the spectral gap of H , whose dependence on the problem it leaves as an open question [1]. The assessment in Section 3.6 tests those two claims against what the construction implies. Section 3.2 fixes the representation and Section 3.3 builds the Hamiltonian from it. Section 3.4 describes the circuits, the filter and the readout the paper uses, and Section 3.5 its route to nonlinear equations. Section 3.6 then judges the method and poses the three questions that Sections 4 to 6 answer.

3.2 The Chebyshev Latent Space Model

The Chebyshev polynomials of the first kind $T_\ell(x)$ are orthogonal on $[-1, 1]$ under the weight $(1 - x^2)^{-1/2}$, and a sufficiently smooth function f on this interval has the convergent expansion $f(x) = \sum_{\ell=0}^{\infty} \hat{f}_\ell T_\ell(x)$ [4]. The model is a truncation of this series at degree $N - 1$, with $N = 2^n$, whose N coefficients are stored as the amplitudes of an n -qubit state. Doubling the truncation

dimension therefore costs one qubit.

The overlap that evaluates the model at a point x is taken with the feature vector

$$|\tau(x)\rangle_n = \frac{1}{2^{n/2}} T_0(x) |0\rangle^{\otimes n} + \frac{1}{2^{(n-1)/2}} \sum_{k=1}^{N-1} T_k(x) |k\rangle. \quad (3.1)$$

The weights $2^{-n/2}$ and $2^{-(n-1)/2}$ normalise the basis: the N vectors $|\tau(\bar{x}_j)\rangle_n$ at the Chebyshev nodes $\bar{x}_j = \cos[\pi(j + \frac{1}{2})/N]$ are orthonormal. At a general x , $|\tau(x)\rangle_n$ is not of unit norm (its norm is below $\sqrt{2}$, since $|T_k| \leq 1$), and this is why the model value is an overlap and not a probability amplitude.

A unit-norm coefficient state $|\psi\rangle_n \in \mathbb{C}^N$ cannot carry the amplitude of the function it represents, so the amplitude is split off as a positive real factor $\sqrt{\eta}$:

$$f_Q(x) = \sqrt{\eta} f_q(x), \quad f_q(x) = \langle \tau(x) | \psi \rangle_n. \quad (3.2)$$

Here f_Q is the physical field and f_q the unit-norm model. Every constraint that follows is a homogeneous linear condition $C|\psi\rangle = 0$, which is unchanged when $|\psi\rangle$ is rescaled, so the Hamiltonian built from such conditions depends on the direction of $|\psi\rangle$ alone. The scale $\sqrt{\eta}$ plays no role until Equation (3.8), which fixes it from one known value of the solution after the ground state has been found. Accordingly, every equation from here until that point is written for $|\psi\rangle$, with η suppressed.

3.3 The Residual, the Hamiltonian and Its Constraints

Each piece of information about the solution is written as a matrix C with $C|\psi\rangle = 0$ when $|\psi\rangle$ holds the coefficients of the solution. What makes this possible is that the components of $\langle \tau(x) |$ are linearly independent functions of x . A differential equation holds at every $x \in [-1, 1]$, so substituting the model gives an overlap $\langle \tau(x) | C |\psi\rangle = 0$ for every x ; since no non-zero combination of the components vanishes identically, this is equivalent to the finite matrix equation $C|\psi\rangle = 0$, and the dependence on x drops out. The one obstacle is a term in which x multiplies the model, such as xf'' : the factor x is not a matrix acting on $|\psi\rangle$. The lifting operators below move it into the feature vector instead.

Differentiating a Chebyshev series produces another Chebyshev series whose coefficients are a fixed linear function of the original ones [1], [4]. The reason this map is dense is the identity $T'_p = 2p \sum T_k$, summed over the $k < p$ of opposite parity to p (the $k = 0$ term has weight p): a single high-degree polynomial differentiates into every lower polynomial of the right parity. In

the normalised basis of Equation (3.1) this is an $N \times N$ matrix $\mathbb{G}$ with

$$[\mathbb{G}]_{k,p} = \begin{cases} 2p, & p > k \geq 1, \text{ } p - k \text{ odd,} \\ \sqrt{2}p, & k = 0, \text{ } p \text{ odd, } p \geq 1, \\ 0, & \text{otherwise,} \end{cases} \quad (3.3)$$

so that $f'_q(x) = \langle \tau(x) | \mathbb{G} | \psi \rangle_n$, and the m -th derivative follows from the power $\mathbb{G}^m$. Here $\mathbb{G}$ acts on coefficient vectors; Wu et al. print the same matrix as $\mathbb{G}_n^\top$. Equation (3.3) is strictly upper triangular, and row k is non-zero at every column $p > k$ of the right parity, so the lowest coefficient of the derivative depends on half of the coefficients of the function. $\mathbb{G}$ therefore carries $\Theta(N^2)$ non-zero entries, and $\|\mathbb{G}\|_2 = \Theta(N^2)$. Appendix B gives the fill-in explicitly at $N = 4$, alongside the banded factors $\tilde{\mathbb{B}}, \tilde{\mathbb{D}}$ that Section 5.2 builds from the same recurrence.

For a homogeneous linear equation of order k with constant coefficients, $\sum_{j=0}^k a_j f^{(j)} = 0$, substituting $\mathbb{G}^j$ for each derivative gives the operator form

$$A_{\text{std}} = \sum_{j=0}^k a_j \mathbb{G}^j, \quad (3.4)$$

which acts on the same n qubits as $|\psi\rangle$. A second-order equation gives $A_{\text{std}} = a_2 \mathbb{G}^2 + a_1 \mathbb{G} + a_0 I$, and the introductory example $f'' + f = 0$ gives $A_{\text{std}} = \mathbb{G}^2 + I$. If $|\psi\rangle$ held the coefficients of an exact solution that is a polynomial of degree below N , then $A_{\text{std}} |\psi\rangle$ would vanish identically.

A coefficient $a_j(x)$ cannot be substituted this way, because $a_j(x)f^{(j)}(x)$ is a product of two functions and the Chebyshev coefficients of a product are not the product of the coefficients. The paper's remedy is a lifting operator $\mathbb{M}_{x^p}$ that carries the factor x^p across the overlap. It rests on the product-to-sum identity

$$x^p T_k(x) = 2^{-p} \sum_{i=0}^p \binom{p}{i} T_{|k+p-2i|}(x), \quad (3.5)$$

the p -fold iterate of the pairwise rule $T_j T_k = \frac{1}{2}(T_{j+k} + T_{|j-k|})$ that reappears below in the nonlinear extension. Conjugating Equation (3.5) into the normalised bases of Equation (3.1) gives an operator satisfying

$$x^p \langle \tau(x) |_n = \langle \tau(x) |_{n+1} \mathbb{M}_{x^p}, \quad (3.6)$$

with entries given in Appendix C. Multiplying a degree- $(N-1)$ series by x^p raises its degree by p , so the product needs more room than the N coefficients available. Taking one extra qubit doubles the register, so $\mathbb{M}_{x^p}$ is a $2N \times N$ matrix and is exact for $p \leq N$, with no truncation. For $p = 0$ the lift reduces to the embedding $\mathbb{M}_1 = \sqrt{2} [I; 0]$, whose factor $\sqrt{2}$ is the ratio of the normalisations of $|\tau\rangle_{n+1}$ and $|\tau\rangle_n$. A coefficient with the finite power series $a_j(x) = \sum_p c_{j,p} x^p$ (a Maclaurin truncation, for a non-polynomial coefficient) lifts by linearity to $\mathbb{M}_{a_j} = \sum_p c_{j,p} \mathbb{M}_{x^p}$,

and the row for the variable-coefficient case is

$$A_{\text{std}} = \sum_{j=0}^k \mathbb{M}_{a_j} \mathbb{G}^j. \quad (3.7)$$

A constant coefficient has $\mathbb{M}_{a_j} = a_j \mathbb{M}_1$, which differs from Equation (3.4) only by the embedding. In the paper's panels an axis is lifted only when it carries a non-constant coefficient or a source: constant-coefficient equations stay on n qubits (its Figs. 3 and 6). When an axis is lifted, every row on that axis is lifted with it, including the data constraints below, so that all rows act into the same $(n+1)$ -qubit register.

The equation alone has a k -dimensional space of solutions, so conditions at chosen points select one. Each condition is built from a rank-one point-evaluation generator, $\mathbb{B}_n(x) = \sqrt{2^n} |0\rangle^{\otimes n} \langle \tau(x)|_n$, which is non-zero only in its first row and satisfies $\langle \tau(x') | \mathbb{B}_n(x) | \psi \rangle = \langle \tau(x) | \psi \rangle$ for every x' . It reads off the model's value at the point x . The paper distinguishes two kinds of condition by how they behave when $|\psi\rangle$ is rescaled.

An *invariant* constraint has a zero target and is unchanged by $|\psi\rangle \rightarrow c|\psi\rangle$. It enters the Hamiltonian directly as one further linear row, $\mathbb{B}|\psi\rangle = 0$: a Dirichlet condition $f(x_z) = 0$ gives $\mathbb{B}_n(x_z)|\psi\rangle = 0$, and a Neumann condition $f'(x_m) = 0$ gives $\mathbb{B}_n(x_m)\mathbb{G}|\psi\rangle = 0$.

A *regular* constraint has a non-zero target, $f(x_s) = y_s$, and is not unchanged by rescaling. It fixes the amplitude of Equation (3.2) directly, after the ground state has been found:

$$\sqrt{\eta} = \frac{y_s}{\langle \tau(x_s) | \psi_g \rangle_n}. \quad (3.8)$$

This is the only place η is determined, and it requires a point x_s at which the solution is non-zero. The same point supplies the *collapse* row $\mathbb{D}^{(0)}(x_s) = \mathbb{B}(x_s)/y_s$, which maps $|\psi\rangle$ to the constant function $f_q(x_s)/y_s$. On the solution this constant equals $1/\sqrt{\eta}$, so $\mathbb{D}^{(0)}$ is a row that tracks the amplitude. A homogeneous linear equation needs only the first kind of condition for the equation itself, so for it the regular datum acts after the ground state has been found. It enters the residual only where the equation needs an amplitude to tie a term to: a source below, and the linear terms of a nonlinear equation in Section 3.5.

The right-hand side r of an inhomogeneous equation $\sum_j a_j f^{(j)} = r$ does not scale with $|\psi\rangle$, so it cannot be a matrix acting on $|\psi\rangle$ alone. The collapse row supplies the missing scaling: $\sqrt{\eta} \mathbb{D}^{(0)}(x_s) |\psi\rangle$ is the constant function 1 on the solution, so multiplying that constant by the series of r gives a term that rescales with $|\psi\rangle$ as the other terms do. A non-polynomial source is expanded as a Maclaurin series truncated at order $\bar{p}$ ($\bar{p} = 3, 5, 7$ in the paper's inhomogeneous panels) and lifted like a coefficient, and the equation row becomes

$$A_{\text{std}} = \sum_{j=0}^k \mathbb{M}_{a_j} \mathbb{G}^j - \mathbb{M}_r \mathbb{D}^{(0)}(x_s), \quad (3.9)$$

where $\mathbb{M}_r = \sum_{p \leq \bar{p}} c_p \mathbb{M}_{x^p}$ is the lift of the truncated series. The truncation is a source of error in its own right, which Section 8.1 measures.

For equations in two variables the state lives on $2n$ qubits, the feature vector is the product of the two one-dimensional ones, and the derivatives act as $\mathbb{G} \otimes I$ and $I \otimes \mathbb{G}$. The heat equation $\partial_t f - \kappa \partial_x^2 f = 0$ gives $A_{\text{std}} = \mathbb{G} \otimes I - \kappa I \otimes \mathbb{G}^2$. The invariant constraints take the form $\mathbb{B}_n(x_z) \otimes I$, and similarly for the second axis. The paper’s partial-differential panels take their extra invariant constraints from zeros of the known solution: $\partial_x f(t, \pm \frac{3}{4}) = 0$ for the heat equation and $f(\frac{1}{8}, x) = f(\frac{7}{8}, x) = 0$ for the wave equation. The initial condition itself is not imposed as a constraint; it enters only through the regular datum, which sets the scale.

Definition 10 (Physics-informed effective Hamiltonian): Collecting the differential-equation row A_{std} and every invariant data constraint $\mathbb{B}_i$,

$$H_{\text{std}} = A_{\text{std}}^\top A_{\text{std}} + \sum_i \mathbb{B}_i^\top \mathbb{B}_i, \quad \langle \psi | H_{\text{std}} | \psi \rangle = \|A_{\text{std}} \psi\|^2 + \sum_i \|\mathbb{B}_i \psi\|^2. \quad (3.10)$$

H_{std} is real, symmetric and positive semi-definite, so $\langle \psi | H_{\text{std}} | \psi \rangle \geq 0$ with equality exactly when every constraint holds simultaneously. The differential equation is thereby an eigenvalue problem, $H_{\text{std}} |\psi_g\rangle = \lambda_{\min} |\psi_g\rangle$ with $\lambda_{\min} \approx 0$, whose ground state $|\psi_g\rangle$ is the solution.

The Hamiltonian is a sum of squared violations: its energy is the total error of the model against the equation and its data. Writing $R = [A_{\text{std}}; \mathbb{B}_1; \dots]$ for the stacked residual, $H_{\text{std}} = R^\top R$, so $|\psi_g\rangle$ is the smallest right singular vector of R and $\lambda_{\min}$ is the square of its smallest singular value. The construction is a least-squares one, not an exact-kernel one. A solution that is a polynomial of degree below N , with every coefficient and source represented exactly, annihilates every row and gives $\lambda_{\min} = 0$; the paper’s Legendre panels of its Figs. 4a and 4b are of this kind, while those with a $\sqrt{1-x^2}$ branch point (Figs. 4c and 4d) are not. For any other solution the truncation leaves a residual, $\lambda_{\min} > 0$ is small, and the ground state is the best fit within the truncated space. A homogeneous equation of order k has a k -dimensional solution space, so $k-1$ invariant constraints leave a one-dimensional (near-)kernel, and the regular datum then fixes its scale; the paper’s second-order panels use one invariant constraint each. The worked $N=4$ residual for the paper’s pilot equation and the entries of the lifts are given in Appendix C.

From here on, the paper’s construction, in which the derivatives are eliminated into the dense matrix $\mathbb{G}$ and the residual is A_{std} with Hamiltonian H_{std} , is called the *standard form*, and the subscript “std” marks its objects. It is the form against which Section 4 and the later chapters compare the reformulation of Section 5, the descriptor form.

3.4 The Published Circuits, Filter and Readout

Solving the equation on a quantum computer needs three subroutines: a block encoding of every term of the Hamiltonian, a filter that isolates its ground state, and a measurement that

reads the field from that state. This section describes each as the paper presents it, with the claims the paper makes for it. The assessment of these claims is in Section 3.6.

Block encodings. The Hamiltonian is a sum of products of terms, so each term is block-encoded separately (Section 2.4). The rank-one rows $\mathbb{B}(x)$ and $\mathbb{D}^{(0)}(x_s)$ are built from the inverse Chebyshev feature map, which prepares $|\tau(x)\rangle$ from the all-zero state, and a reflection $\hat{S}_{n+1} = 2|0\rangle\langle 0| - I$ about the all-zero state, combined by a two-term linear combination of unitaries (LCU) [5]. The encoding of $\mathbb{B}$ uses $n + 2$ qubits with subnormalisation $\sqrt{2^{n+1}}$, and that of $\mathbb{D}^{(0)}(x_s)$ divides it by $f(x_s)$. The dense $\mathbb{G}$ is also block-encoded as an LCU, which exploits the structure of Equation (3.3): the entries grow linearly along each superdiagonal, and two conditioned subroutines load them as rotation amplitudes, with a small error in each entry [1]. The circuit’s unitary contains $\mathbb{G}$ in its top-left block, divided by the prefactor

$$\alpha_{\mathbb{G}}^{\text{pub}} = (2^{n-1} + 2^n) \|\mathbb{G}\|_S, \quad \|\mathbb{G}\|_S = \max(\|\mathbb{G}\mathbb{G}^\top\|_1, \|\mathbb{G}^\top\mathbb{G}\|_1). \quad (3.11)$$

The paper’s claims for this encoding are that it needs $2n + 3$ qubits; that its gate count is linear in n for each of $\mathbb{D}^{(0)}(x_s)$ and $\mathbb{G}$, and $O(n^3)$ overall once the multi-controlled gates are decomposed; and that the entry error decreases exponentially in n , for which no bound is given. The registers, subroutines and angles are given in Appendices D.1 and D.2, and the claims are tested in Section 4.1.

The imaginary-time filter. The ground state is prepared by quantum imaginary-time evolution (quantum imaginary-time evolution (QITE)): applying $e^{-\beta H}$ to an initial state multiplies each eigenvector by $e^{-\beta E}$ for its energy E , so that the lowest survives once β exceeds the inverse gap. The paper realises this by quantum singular value transformation (QSVT) [6] on the block-encoded $\tilde{H} = H/\|H\|_F$, with $\|\cdot\|_F$ the Frobenius norm and with phase angles fitted to $e^{-\beta x}$ (Sections 2.5 and 2.6 give the machinery). The evolution time of $e^{-\beta \tilde{H}}$ is bounded below by

$$\beta_{\text{pub}} \geq \frac{\lambda_{\max}}{\lambda_2(n-1)}, \quad (3.12)$$

where $\lambda_{\max}$ and λ_2 are the largest and the second smallest eigenvalues (the paper writes t for β , calls the bound approximate, and states it for H or $\tilde{H}$ alike, the ratio being scale-free). Since $\lambda_{\min} \approx 0$, λ_2 plays the role of the gap Δ . The paper expects the preparation time to scale as $O(1/\Delta) \text{polylog}(1/\varepsilon)$ in the gap Δ , and as $O(1/\sqrt{\Delta})$ with algorithmic thermalisation, and states that the gap’s dependence on the problem is “problem-dependent and in general are not easy to bound” [1]. The phase sets, the dilation of $\tilde{H}$ and the parameters of the paper’s examples are in Appendix K.5.

Readout. The field is read out by an interferometric overlap measurement [7], [8]. The model value $f_Q(x)$ is a signed overlap, but a measurement of the all-zero outcome after the inverse feature map returns only its square. The paper therefore measures that probability

twice: once for the prepared state, and once for that state superposed with a reference state. The two probabilities determine the signed overlap $\langle \tau(x) | \psi_g \rangle$, hence $f_Q(x)$ at the chosen point, with η from the regular datum (Equation (3.8)); where the solution has a single sign, the first probability alone suffices (Eqs. 36–39 of the paper). The field is therefore reconstructed point by point. The identity, in the form this thesis uses, is Equation (7.12).

3.5 The Doubled-Space Extension to Nonlinearity

A nonlinear term such as $(f')^2$ is not linear in $|\psi\rangle$, so no matrix acting on $|\psi\rangle$ can represent it. The paper makes it linear by doubling the Hilbert space: the unknown becomes $|\Psi\rangle = |\psi\rangle \otimes |\psi\rangle$, and a product of two model functions is the overlap of $|\Psi\rangle$ with a product of two feature vectors [1]. The product-to-sum rule $T_j T_k = \frac{1}{2}(T_{j+k} + T_{|j-k|})$ folds this product back into a single Chebyshev series on one extra qubit,

$$\langle \tau(x) |_n \otimes \langle \tau(x) |_n = \langle \tau(x) |_{n+1} \mathfrak{m}_1, \quad \mathfrak{m}_1 \in \mathbb{R}^{2N \times N^2}, \quad (3.13)$$

so a quadratic term $f^{(i)} f^{(j)}$ becomes the operator $\mathfrak{m}_1(\mathbb{G}^i \otimes \mathbb{G}^j)$ acting on $|\Psi\rangle$. The linear terms must act on the same space. They are carried to it by tensoring them with the collapse row $\mathbb{D}^{(0)}(x_s)$, which supplies the constant factor of one copy of the model, and the scale of the doubled model is then η and not $\sqrt{\eta}$. The Hamiltonian is built from the lifted residual and the data constraints as in Definition 10, on $2n$ qubits. The paper demonstrates it on two quadratic ordinary differential equations at $n = 2$.

The construction models a *quadratic* nonlinearity only. A cubic term would need three copies, $|\psi\rangle^{\otimes 3}$, folded back by two applications of the product rule, and in general a term of degree p needs p copies and a register of dimension N^p ; Paine’s concurrent work makes the same observation [9]. Even for quadratic terms, $|\Psi\rangle$ lives in a space of dimension $N^2 = 2^{2n}$, and the paper reports that the Hamiltonian has *at least* 2^{2n-1} degenerate zero-energy states, half of that space at minimum. Solving the equation is, in the paper’s words, “converted to that of seeking a degenerate state” that “closely matches the analytical solution” [1]; the text gives no rule for that selection that does not refer to the known solution.

Assessment: projection, not preparation. The degeneracy matters for the filter. Any filter that is a function of the Hamiltonian acts as a scalar on its zero-energy eigenspace, so it returns the normalised projection of the initial state onto the kernel, and the direction chosen inside the kernel is set by the initial state and not by the equation. Nor can the Hamiltonian be adjusted to single out the solution. A kernel is a subspace, so one that contains $|\psi\rangle \otimes |\psi\rangle$ and $|\phi\rangle \otimes |\phi\rangle$ also contains their sum, which is not a product state unless $|\psi\rangle$ and $|\phi\rangle$ are parallel. No penalty built without knowledge of the solution can therefore make the product structure part of the kernel. Extracting the physical solution from the doubled space is selection within a known kernel and not preparation of a state. The formal statement, with its scope, is Proposition 4, and Section 6 constructs, for a restricted but practically important class of

equations, a preparation that does not have this property.

3.6 Critical Assessment

Strengths. The construction is unified: ordinary and partial differential equations, and the data constraints, take the same least-squares form, and a nonlinear term has a route into it. The kernel is exact whenever the solution is a polynomial within the truncation. The Chebyshev basis needs no spatial grid, so the model and its derivatives are available at every $x \in [-1, 1]$ from the one prepared state. The paper reports that each term of the Hamiltonian block-encodes with a number of gates polynomial in n [1]. These are real advantages over grid-based formulations, and the framework is worth building on.

Weaknesses. What follows ranks the weaknesses by how much each limits the method’s claim to efficiency, most severe first. Each is tested later, and the pointer names where. The first two concern the cost, the next three the evidence of the paper’s panels, and the last two the extension to nonlinearity and the readout.

1. **The spectrum fixes the cost.** For an equation of order k , $\|\mathbb{G}\|_2 = \Theta(N^2)$ gives $\|H_{\text{std}}\| = \Theta(N^{4k})$, which is N^8 at second order; the exponent doubles once because the Hamiltonian is a Gram product, $\|H_{\text{std}}\| = \|R\|^2$, the penalty of normal equations. The filter degree is $d = \Theta(\sqrt{\alpha_H/\Delta} \log 1/\varepsilon)$, where $\alpha_H \geq \|H_{\text{std}}\|$ for any block encoding. Where the gap Δ does not grow with N , which is the case for the standard form on the one-axis problems measured (Section 4.4), the degree is therefore at least $\tilde{\Theta}(N^{2k})$, which is $\tilde{\Theta}(N^4)$ at second order, *whatever the block encoding*. The method is not polylogarithmic in N , and neither is the reformulation of Section 5. The paper’s expectation $O(1/\Delta) \text{polylog}(1/\varepsilon)$ is correct if Δ is the gap of the normalised Hamiltonian $\tilde{H}$, which is Δ/α_H in the units of the Hamiltonian itself; the cost is then $O(\alpha_H/\Delta)$, and the norm re-enters through the relative gap. The paper does not say in which units Δ is meant. Section 8.3 measures the resulting scaling.
2. **The published normalisation is loose.** The prefactor $\alpha_{\mathbb{G}}^{\text{pub}}$ of Equation (3.11) exceeds $\|\mathbb{G}\|_2$ by a factor that grows as N^3 , so the published encoding is loose even against the matrix it encodes. The first weakness holds for any block encoding; this one is specific to the published circuit and is a cost on top of it (Sections 4.1 and 8.3).
3. **Printed choices distort the solution.** The printed constraint points are rounded, and the inhomogeneous panels use Maclaurin-truncated sources; each sets an error floor at the printed n that no resolution removes (Section 8.1).
4. **The partial-differential panels do not extend.** Their extra invariant constraints are read off the known solution, which a new problem does not supply, and they degenerate beyond the printed n . The heat panel imposes its data at $t = 0$ on a domain that ends there, which makes it a backward and therefore ill-posed problem (Section 8.1).

5. **The filter’s inputs.** The time rule Equation (3.12), which the paper calls approximate, needs the largest eigenvalue and the gap in advance, and the method does not supply them. Tested, it falls short of the target accuracy (Section 8.1). The post-selected state is the ground state only to the extent that the initial state overlaps it, and the success probability falls with that overlap.
6. **The nonlinear route projects.** The doubled space leaves a kernel of dimension at least 2^{2n-1} , so a filter returns the projection of the initial state and cannot prepare the solution; the construction is also quadratic only (Sections 3.5 and 8.5).
7. **The readout assumes a deterministic preparation.** The identity of the paper treats the prepared state as exactly the ground state, although a filter followed by post-selection succeeds with a probability below one. It also returns one point per measurement campaign (Sections 7.7 and 8.1).

The paper’s open problems. The Discussion of the paper names two: “studying the spectral gap dependence is an important question for the future work”, and the choice of basis, for which “one guiding principle here can be the cost of differentiation, where the goal is minimizing the depth of quantum circuits to implement derivatives, multiplication etc.” [1]. Paine’s concurrent work [9] compares alternative Hamiltonian constructions and function encodings and finds, on a few small instances, that the normalised gap depends on both. It states that these results do not establish a general ordering, and it gives neither the gap’s scaling with resolution nor block-encoding circuits. It raises the first open problem without answering it. The first weakness above is the quantitative form of that problem: the gap is a property of the Hamiltonian, and the norm that it must be set against is fixed by the differentiation matrix.

Remaining questions. Three questions follow, each answered by one chapter. *(i) Is the cost of the method inherent to the Hamiltonian that its construction defines, or is it incurred by the particular circuits used to block-encode that Hamiltonian?* Section 4 rebuilds the published block encoding and builds an exact structured one beside it, so that the cost of the circuit can be separated from the cost set by the spectrum. *(ii) Does a reformulation of the residual exist whose Hamiltonian is banded and well conditioned and whose solution is unchanged, and what does it cost in register size and spectral gap?* Section 5 carries the derivatives as unknowns tied by banded recurrences, instead of eliminating them into the dense $\mathbb{G}$. *(iii) Can the solution of a nonlinear equation be prepared as the unique ground state of a Hamiltonian, rather than selected by projection from a degenerate kernel?* Section 6 sets the doubled space against a Carleman lift, which prepares a state for a restricted class of polynomial equations, and measures separately whether the lift converges. Underlying all three is a single question: whether the norm of the Hamiltonian, and with it the filter’s degree, can be lowered by changing what the residual carries, and at what price.

4 The Standard Form: Encoder and Hamiltonian Costs

This chapter answers the first of the three questions of Section 3.6: is the cost of the physics-informed effective Hamiltonian the cost of the circuit that block-encodes it, or the cost of the Hamiltonian itself? We first rebuild the published encoding of the differentiation matrix and test the paper’s claims for it (Section 4.1). We then construct an exact encoding of the same matrix whose subnormalisation is within a constant of its norm (Sections 4.2 and 4.3), so that for constant coefficients the cost that survives it is not the encoder’s (Section 4.4). This structured encoding is a fairness device: it gives the comparison of Section 8.3 the strongest baseline we can build for the paper’s regime, and it is not offered as a solver for that regime.

4.1 The Published Encoding, Rebuilt

Wu et al. report that each term of the Hamiltonian block-encodes with a number of gates polynomial in n (Section 3.4). We rebuild their three circuits, Figs 2, 8 and 9 of [1], and test each claim against the build; the registers, angles and every deviation from the printed figures are in Appendix D.1. The zero row $\mathbb{B}(x)$ and the collapse row $\mathbb{D}^{(0)}(x_s)$ (Figs 2 and 8) rebuild as printed, with $n + 2$ qubits and subnormalisation $\sqrt{2^{n+1}}$ (divided by $|f(x_s)|$ for $\mathbb{D}^{(0)}$, the sign carried as a phase). They succeed with probability $\|\tau(x)\|^2/2$, one half on average over the Chebyshev measure, so the paper’s “high probability” is a constant, not certainty.

The derivative circuit (Fig. 9) cannot be rebuilt wire for wire, because its multi-controlled wiring is not legible in the figure. We rebuild its mechanism with registers of our own: a uniform superposition over the $N/2$ odd offsets $p-k$, a register subtraction that selects the entry (k, p) , and ladders of small controlled rotations that load the column value $2p$ as an amplitude and correct row 0. The ladder composes angles exactly, but $\sin(\sum_m \arcsin(2^m \mu)) \neq \mu p$, so the loaded entries are approximate. The block differs from $\mathbb{G}$ entrywise by at most $\frac{1}{3}\mu^2 p^3$ to leading order (Proposition 8), which at the printed $\mu = 2/\|\mathbb{G}\|_S$ is $O(N^{-5})$ and falls by a factor approaching 32 per qubit.

Table 2 gives the verdict on each claim. Four hold, the qubit count is two higher, and the prefactor is valid but loose; the last decides the cost. A circuit with polynomially many gates is cheap to apply, but the number of times the filter must apply it is set by the subnormalisation (Equation (2.6)), and the printed prefactor Equation (3.11) exceeds $\|\mathbb{G}\|_2$ by a factor of 2.2×10^6 at $n = 7$. That factor tends to jN^3 , where $j \approx 1.06$ is the first positive zero of the Bessel function $J_{-3/4}$ (Appendix E.4). It is the first of the two walls this chapter separates, the encoder’s: the paper’s claim that the terms are implementable efficiently is true of the gates and false of the subnormalisation. Our build’s own α is a third of the printed prefactor ($N\|\mathbb{G}\|_S/2$ against $3N\|\mathbb{G}\|_S/2$), so every published-encoding cost in this thesis is the kinder side of the paper’s. The encoding is approximate, so no state is prepared with it; it enters only the resource estimates,

composed into the same stacked residual as the structured encoding, so that its filter degree is counted like the others (Sections 4.4 and 8.3).

Table 2: The paper’s claims for its block encodings, tested against the rebuilt circuits. Figs 2 and 8 hold as printed. In Fig. 9, the gate count and the error hold, the qubit count is two higher, and the prefactor is valid but loose against $\|\mathbb{G}\|_2$ by a factor that grows as N^3 (Table 3); the polynomial gate count does not make the encoding efficient. Classical exact-mode build; circuit counts as built, error to leading order.

Circuit and claim	Printed	As rebuilt	Verdict
Fig. 2, $\mathbb{B}(x)$	$n + 2$ qubits, $\sqrt{2^{n+1}}$	the same, exact	holds
Fig. 8, $\mathbb{D}^{(0)}(x_s)$	$\sqrt{2^{n+1}}/f(x_s)$	the same, exact	holds
Fig. 9, qubits	$2n + 3$	$2n + 5$	two more
Fig. 9, gates	$O(n^3)$	$O(n^2)$	holds (as a bound)
Fig. 9, error	decreases exponentially	$O(N^{-5})$, bound derived	holds
Fig. 9, prefactor	$(2^{n-1} + 2^n)\ \mathbb{G}\ _S$	valid, $\sim jN^3\ \mathbb{G}\ _2$	loose

4.2 An Exact Factorisation of the Differentiation Matrix

The published circuit loads the column value as a rotation amplitude, which is where it is approximate and where its prefactor comes from. The structure of $\mathbb{G}$ allows a different route. Every non-zero entry is a product of three factors, each depending on one index alone: a weight depending only on the row, smaller by $1/\sqrt{2}$ in row 0 because of the normalisation of T_0 ; a selector depending only on the offset $p - k$, non-zero when $p > k$ and $p - k$ is odd; and a value depending only on the column, $2p$, the factor by which differentiation weights the lower terms of T_p . Each factor has an exact block encoding.

Proposition 1 (Exact factorisation of the differentiation matrix). *For every $n \geq 1$ the differentiation matrix on the normalised Chebyshev coefficients factorises as*

$$\mathbb{G} = R_0 (2U_{\text{odd}}) \Lambda, \quad (4.1)$$

$$R_0 = I - \left(1 - \frac{1}{\sqrt{2}}\right) |0\rangle \langle 0|, \quad (U_{\text{odd}})_{kp} = [p > k, p - k \text{ odd}], \quad \Lambda = \text{diag}(p).$$

Here U_{odd} is the sum of the $N/2$ truncated shifts by odd offsets.

The proof is entry by entry against Equation (3.3) and is given in Appendix E.1. The identity is exact. As a check on our implementation and not on the proposition, the composed circuit’s block matches the independently constructed $\mathbb{G}$ entrywise to 3.4×10^{-13} (absolute) for $n \leq 7$. The factorisation also sets the size of $\mathbb{G}$. The subnormalisation $N(N - 1)$ of Proposition 2 bounds $\|\mathbb{G}\|_2$ above by $O(N^2)$, and the matching lower bound, with $\|\mathbb{G}\|_2 \rightarrow N^2/(2j)$ for the j above, is proved in Appendix E.3.

4.3 Circuits and Cost

Each factor is encoded as a linear combination of unitaries (LCU) [5], with no value loaded by rotation, so none of the encoding is approximate. The column index Λ is a sum of Pauli strings,

$$\Lambda = \frac{N-1}{2} I - \sum_{j=0}^{n-1} 2^{j-1} Z_j, \quad (4.2)$$

which follows from $p = \sum_j 2^j(1 - z_j)/2$ for the eigenvalue z_j of Z_j . It has $n+1$ terms, so $\lceil \log_2(n+1) \rceil$ selector qubits and $O(n)$ gates, and its subnormalisation $\alpha_\Lambda = N-1$ equals $\|\Lambda\|_2$. The row weight is $R_0 = aI + b(I - 2|0\rangle\langle 0|)$ with $a = \frac{1}{2}(1 + 1/\sqrt{2})$ and $b = \frac{1}{2}(1 - 1/\sqrt{2})$, a two-term LCU with $\alpha_{R_0} = 1 = \|R_0\|_2$.

The sum of shifts is the one non-trivial factor (Figure 1). Its PREPARE places the offset register in a uniform superposition over the $N/2$ odd offsets, by an X on the lowest qubit and Hadamards on the rest. Its SELECT is the in-place subtraction $p \leftarrow p - d$ of a ripple-carry circuit [10], which also sets a borrow flag when $p < d$. The terms that would wrap around carry a borrow of one, and post-selecting the flag on zero removes them exactly. The block is therefore $\frac{2}{N}[p > k, p - k \text{ odd}]$, an exact encoding of U_{odd} at $\alpha = N/2$ with $O(n)$ Toffolis, and a single PREPARE serves all $N/2$ terms because SELECT is a permutation (Appendix E.2).

Quantikz circuit: offset register d (X on the lowest qubit, $H^{\otimes(n-1)}$ on the rest), system register p , carry and borrow ancilla; a ripple-carry subtractor $p \leftarrow p - d$ spanning all five wires; $\text{PREPARE}^\dagger$; post-selection of d and the borrow flag on $|0\rangle$.

Figure 1: **The odd-offset shifts are one subtraction.** PREPARE (left) puts the offset register d in a uniform superposition over the $N/2$ odd offsets, by an X on its lowest qubit and Hadamards on the rest. SELECT is a ripple-carry subtractor [10] that replaces the system register p by $p - d$ and sets the borrow flag when $p < d$. $\text{PREPARE}^\dagger$ (right) with post-selection of d and the borrow flag on $|0\rangle$ keeps $\frac{2}{N}|k\rangle\langle p|$ for every odd $d = p - k \geq 1$, and the wrapped terms vanish exactly. The block is $U_{\text{odd}}/(N/2)$, so $\alpha = N/2$, on n system and $n+2$ ancilla qubits.

The three encodings multiply, and their subnormalisations multiply with them. The composition is an exact block encoding of $\mathbb{G}$ with $\alpha_{\mathbb{G}} = 1 \cdot N \cdot (N-1)$, and the cost is the following.

Proposition 2 (Cost of the structured encoding). *The factorisation of Proposition 1 compiles to an exact block encoding of $\mathbb{G}$ on n system qubits with*

$$\alpha_{\mathbb{G}} = N(N-1), \quad \frac{\alpha_{\mathbb{G}}}{\|\mathbb{G}\|_2} \longrightarrow 2j_{-3/4,1} \approx 2.117, \quad O(n^2) \text{ gates}, \quad n+4 \text{ ancilla}. \quad (4.3)$$

Of the ancilla, $n - 1$ are the selector of U_{odd} 's $N/2$ -term LCU.

The limit is the Bessel-kernel limit of $\|\mathbb{G}\|_2/N^2$ already used above, the gate count is for an ancilla-free synthesis of multi-controlled gates and is $O(n \log n)$ with a linear-cost synthesis that borrows idle qubits [11], and the proof is given in Appendix E.3. The ancilla count is exact in the build for $n = 2$ to 8, and the gate counts are the transpiled two-qubit counts at small n only (109 at $n = 2$ and 471 at $n = 5$), which are consistent with the bound but too few to fit it. The limit is a statement about large N , and the sizes we reach are measured: the ratio is 1.62 at $n = 2$, rises with n , and is 2.11 at $n = 10$ (Table 3), against a published ratio that grows as N^3 . The looseness is small because α_{R_0} and α_Λ equal their norms; it comes from U_{odd} , whose α exceeds its norm by at most a factor of $\pi/2$, and from the product, whose large factors meet each other's small ones.

The $n - 1$ selector ancilla are not an artefact of this circuit. An LCU of L terms needs $\lceil \log_2 L \rceil$ selector qubits, and lower bounds of this kind hold for exact block encodings built from L -term decompositions [12], so the count is the price of encoding the $N/2$ shifts of U_{odd} as one LCU. It is a difference between the two regimes: the descriptor form of Section 5.7 never forms the sum of shifts, and its ancilla grow as $\log n$.

Table 3: Subnormalisation of the differentiation-matrix encodings, relative to $\|\mathbb{G}\|_2$. The structured encoding of Proposition 2 stays within a factor 2.12 of the norm, and the published prefactor of Equation (3.11) grows as N^3 . The full table to $n = 10$ is Table 14. Classical exact-mode computation.

n	$\ \mathbb{G}\ _2$	Structured (Proposition 2)	Published (Equation (3.11))
2	7.39	1.623	49.9
4	121	1.986	4.05×10^3
7	7739	2.100	2.20×10^6
10	495 308	2.115	1.14×10^9

The remaining terms of the residual use the same pieces. A coefficient $a(x)$ multiplies a series by the exact lift $\mathbb{M}_a$ to $n + 1$ qubits, which holds the product with no truncation whenever $\deg a \leq N$ (Appendix C.2); it is encoded as a sum of powers of one single- x block, itself built from truncated shifts by the same post-selected subtraction, so its subnormalisation is the sum of the absolute monomial coefficients weighted by powers of that block's. The rank-one rows $\mathbb{B}(x)$ and $\mathbb{D}^{(0)}(x_s)$ are the exact circuits of Section 4.1, and a source given by its Chebyshev coefficients is loaded by one rotation tree. A power $\mathbb{G}^k$ is a product of encodings that share one ancilla register with a failure flag per product, so $\mathbb{G}^2$ costs one qubit more than $\mathbb{G}$; the subnormalisation of a product is the product of the subnormalisations, so a power compounds the constant above while $\|\mathbb{G}^k\|_2 < \|\mathbb{G}\|_2^k$. The stacked residual is then assembled as a reflection (Section 7.4), and its composed circuit is verified against the independently assembled Hamiltonian (Section 7.6).

4.4 What Remains: The Standard-Form Hamiltonian

For constant coefficients the encoder has done its job. On R2a at $n = 8$ the structured encoding of the whole stacked residual has $\alpha_R = 4.26 \times 10^9$ against $\|R\|_2 = 5.51 \times 10^8$ (a resource estimate, not a simulation), a ratio $\alpha_R/\|R\|_2$ of 7.74. The ratio rises from 4.72 at $n = 2$, with increments that shrink by about half at each added qubit, so it is levelling off over $n = 2$ to 8; we have not proved a limit for the whole residual (Section 8.3). The Hamiltonian is encoded at $\alpha_H = \alpha_R^2$, so $\alpha_H/\|H_{\text{std}}\|$ is that ratio squared, from 22.3 to 59.9. The published encoding's ratio grows without such a limit, as N^3 for $\mathbb{G}$ alone.

What the encoder cannot change is the norm it encodes. For an equation of order k with constant coefficients, a non-zero leading coefficient and constraints of lower order, $\|\mathbb{G}^k\|_2/N^{2k}$ tends to a positive constant, so $\|R_{\text{std}}\|_2 = \Theta(N^{2k})$ and $\|H_{\text{std}}\| = \|R_{\text{std}}\|_2^2 = \Theta(N^{4k})$, the proof being given in Appendix E.3. The factor N^2 per derivative is $\mathbb{G}$'s own norm (Section 4.2), and the exponent doubles once more because the Hamiltonian is the Gram product $R^\dagger R$. The gap, in contrast, is not proved flat. We have measured it on R2a and R2b, the one-axis panels with a non-vanishing leading coefficient checked so far: on R2a it is $\Delta = 7.04$ at $n = 7$ and 7.04 at $n = 8$, while $\|H_{\text{std}}\|$ grows from 1.2×10^{15} to 3.0×10^{17} . Since $\alpha_H \geq \|H_{\text{std}}\|$ for any block encoding, Equation (2.6) gives

$$\|H_{\text{std}}\| = \Theta(N^{4k}), \quad \Delta = \Theta(1) \text{ (measured)} \implies d = \tilde{\Theta}(N^{2k}), \quad (4.4)$$

which is $\|H_{\text{std}}\| = \Theta(N^8)$ and $d = \tilde{\Theta}(N^4)$ at second order, the order of the paper's panels, whatever the block encoding. We call this the Hamiltonian wall. The floor is stated for the minimax filter that prepares every state in this thesis; for polynomial filters in general it is what the optimality of Lin and Tong's construction supports [13], and we claim no more. At $n = 8$ the estimated degree for R2a is 1.2×10^{10} as built and 1.5×10^9 at the ideal-encoding bound $\alpha = \|R\|_2$. The degree is linear in α_R , so the difference is the encoder's looseness and the common part is the Hamiltonian's. Both are resource estimates, not simulations, since the degrees are far beyond exact emulation.

The two walls show in the growth of the degree with N between $n = 7$ and $n = 8$, again as resource estimates. The published encoding's degree grows as $N^{10.0}$ and the structured encoding's as $N^{4.01}$, so the encoder wall is the difference between them. At the ideal bound the degree grows as $N^{4.00}$ however the matrix is encoded, and that exponent is the Hamiltonian wall. It belongs to the least-squares functional, not to the circuits. The physics-informed Hamiltonian $H_{\text{std}} = R^\dagger R$ is the normal-equations form of the least-squares problem for an order- k operator, and such a problem squares the conditioning of the operator it minimises. The classical remedies are to recast the equation as a first-order system [14] or to use a basis in which differentiation is banded (Sections 2.3.2 and 2.3.3).

The scope of this conclusion has a limit. The structured encoding bounds $\alpha_H/\|H_{\text{std}}\|$ for constant coefficients only, and it is loose for a variable coefficient that is lifted. On R3, whose

leading coefficient is $1 - x^2$, the circuit pays for the product $\mathbb{M}_a \mathbb{G}^2$ as $\|\mathbb{M}_a\| \|\mathbb{G}^2\| \gg \|\mathbb{M}_a \mathbb{G}^2\|$, because the coefficient vanishes at the ends of the interval where $\mathbb{G}^2$ is largest. The measured norm of R3's residual then grows as N^2 and not as $N^{2k} = N^4$, and the encoder wall reappears: the as-built degree grows as $N^{4.01}$ while the degree at the ideal bound grows as $N^{2.00}$, the exponent that the descriptor's degree also has there ($N^{2.00}$ at the ideal bound). On R3a at $n = 8$ the as-built degree is 1.7×10^{10} against 1.1×10^5 at the ideal bound, both resource estimates. For this reason Section 8.3 quotes the ideal-encoding bound beside every as-built cost, and where the comparison between the forms reverses at the bound, as it does on R3, the reversal is reported as a finding.

This answers the first question. For the constant-coefficient problems of the paper's regime, the cost of the published method has two parts. One is the encoder's, a factor that grows as N^3 in the subnormalisation and is removed by the structured encoding. The other is the Hamiltonian's, $\tilde{\Theta}(N^{2k})$ in the filter degree, which no encoder removes. Section 8.3 tests this on the benchmark problems. What can still change is the Hamiltonian itself, by changing what its residual carries. Section 5 does so by carrying the derivatives as unknowns instead of eliminating them into the dense $\mathbb{G}$.

5 The Descriptor Reformulation

This chapter answers the second question of Section 3.6: can the residual be reformulated so that its Hamiltonian is banded and well conditioned without changing its solution, and what does that cost? We build the reformulated residual and prove that its kernel is the standard form's (Sections 5.1 to 5.4), then state its own challenge, an eroded spectral gap, and the cure with its limits (Sections 5.5 and 5.6), and finally its block encoding, cost and consequence for the filter's degree (Sections 5.7 to 5.9).

5.1 The Idea

Section 4 located the cost of the standard form in its Hamiltonian: $\|\mathbb{G}^k\|_2 = \Theta(N^{2k})$, so $\|H_{\text{std}}\| = \Theta(N^{4k})$ whatever the encoder. That norm is the density of $\mathbb{G}$ made visible: a dense upper-triangular operator is a sum of $N/2$ shifted diagonals (Section 4.3), a banded one a fixed number of them. The density is not intrinsic to differentiation. It comes from insisting that the derivative's coefficients be expressed in the same basis as the function's. If the derivative may live in a neighbouring basis, differentiation is banded, and the factors of $\mathbb{G}$ can be kept instead of multiplied out.

The descriptor form does this. Where the standard form substitutes $\mathbb{G}^j$ for each derivative, the descriptor form carries $f', \dots, f^{(k)}$ as unknowns beside f , ties them to one another by banded recurrences, and writes the equation over the unknowns. A first-order equation or system needs no extra unknowns, because the equation already gives the derivative: multiplying the standard rows through by a banded factor is enough. The first layout we call the chain and the second the mass form (Section 5.3).

Neither ingredient is new to numerical analysis, and we credit both. Banded differentiation is the ultraspherical spectral method of Olver and Townsend [15], of which the factors of Section 5.2 are the discrete shadow (Section 2.3.2). Carrying derivatives as unknowns to avoid squaring the condition number of a least-squares functional is the idea of first-order-system least squares [14] (Section 2.3.3). The physics-informed Hamiltonian is a least-squares functional, so both apply to it. The increment of this thesis is their transfer to a block-encoded physics-informed Hamiltonian: a proof that its kernel is unchanged, an encoding as one flat linear combination of unitaries, an a priori rescaling that restores the gap the new unknowns erode, and a measurement of what all of it does to ground-state preparation.

Figure 2 shows the effect on the second-order equation $f'' + 4f' + 4f = 0$ (R2a). At $n = 7$ the standard form's stacked residual holds 8384 non-zero entries on 128 unknowns; the descriptor's holds 1402 on 512 unknowns, four times as many, because of the block register that Section 5.3 counts. The residual is sparser although its state is larger, and that trade is the subject of the

chapter. The idea in one sentence: carrying the intermediate derivatives instead of eliminating them turns the physics-informed residual from dense to banded with its kernel unchanged, so that the Hamiltonian's norm falls from $\Theta(N^{4k})$, which is $\Theta(N^8)$ at second order, to $\Theta(N^2)$ at fixed block weights; the price is a block register and a gap that the derivative blocks erode, which an a priori rescaling restores on one axis at no gate cost.

Two sparsity (spy) plots side by side for $f'' + 4f' + 4f = 0$ (R2a) at $n = 4$: left, $A_{\text{std}} = \mathbb{G}^2 + 4\mathbb{G} + 4I$ with its dense upper triangle; right, A_{desc} with its block-banded pattern (recurrence blocks, equation row, padding). The descriptor's matrix is four times the size and holds fewer non-zero entries at $n = 7$ (Section 5.1).

Figure 2: Dense against banded: the two residuals of the same equation.

5.2 The Banded Factorisation

Differentiation of a Chebyshev series is a sparse operation written densely. Let $f = \sum_{\ell} a_{\ell} T_{\ell}$ have derivative $f' = \sum_{\ell} p_{\ell} T_{\ell}$. The two coefficient sequences are tied by the recurrence

$$c_{\ell-1} p_{\ell-1} - p_{\ell+1} = 2\ell a_{\ell}, \quad \ell \geq 1, \quad (5.1)$$

with $c_0 = 2$ and $c_{\ell} = 1$ for $\ell \geq 1$ (the factor on p_0 is a normalisation artefact of T_0 , derived with the rest in Appendix B). Each row relates one coefficient of the function to just two of the derivative. Truncated to N coefficients and written in the normalised basis of Section 3.2, where $\hat{f}^{(0)}$ and $\hat{f}^{(1)}$ hold the coefficients of f and f' , this is the matrix identity

$$\tilde{\mathbb{B}} \hat{f}^{(1)} = \tilde{\mathbb{D}} \hat{f}^{(0)}. \quad (5.2)$$

$\tilde{\mathbb{D}}$ has a single non-zero diagonal, the first superdiagonal, with entries $(2\sqrt{2}, 4, 6, \dots, 2(N-1))$; $\tilde{\mathbb{B}}$ has the diagonal $(2, 1, 1, \dots)$ and the second superdiagonal $(-\sqrt{2}, -1, -1, \dots)$. Both hold $\mathcal{O}(N)$ non-zero entries, against the $\Theta(N^2)$ of $\mathbb{G}$. Two facts about them are used throughout. $\tilde{\mathbb{B}}$ is upper triangular with a non-zero diagonal, so it is invertible; and, since the singular values of a single-diagonal matrix are its entries, $\|\tilde{\mathbb{D}}\|_2 = 2(N-1)$ while $\|\tilde{\mathbb{B}}\|_2 \leq 2 + \sqrt{2}$.

The dense differentiation matrix is the quotient

$$\mathbb{G} = \tilde{\mathbb{B}}^{-1} \tilde{\mathbb{D}}, \quad (5.3)$$

an identity that holds exactly at every N . The inverse $\tilde{\mathbb{B}}^{-1}$ is a triangle of ones on each parity class, and the density of $\mathbb{G}$ is that inverse; Equation (5.2) is the same linear map before the inverse is taken. This is the discrete form of the ultraspherical spectral method [15]: up to the basis normalisation, $\tilde{\mathbb{B}}$ and $\tilde{\mathbb{D}}$ are twice its operators for conversion from Chebyshev to second-kind polynomials and for differentiation into them, so that differentiation becomes sparse once the derivative is allowed to sit in a polynomial basis one order above the function (Appendix B). The method applies these operators to the equation to obtain a banded system in f alone. We instead keep Equation (5.2) as a row of the system and the derivative as an unknown, which is what lets the same construction serve every order, every axis and every coupled field.

5.3 Two Layouts and Kernel Equivalence

Consider the linear ordinary differential equation (ODE) of order k

$$\sum_{j=0}^k a_j(x) f^{(j)}(x) = r(x), \quad x \in [-1, 1], \quad (5.4)$$

with the data constraints of Section 3.3. We represent every function in it by its first N coefficients in the normalised Chebyshev basis. Two layouts carry the equation, according to its order.

The chain ($k \geq 2$). Let $\hat{f}^{(j)} \in \mathbb{C}^N$ be the coefficient vector of $f^{(j)}$ for $j = 0, \dots, k$, and stack them into the state

$$w = (\hat{f}^{(0)}; \hat{f}^{(1)}; \dots; \hat{f}^{(k)}). \quad (5.5)$$

Two kinds of row couple the blocks. The *recurrence rows*

$$-\tilde{\mathbb{D}} \hat{f}^{(j)} + \tilde{\mathbb{B}} \hat{f}^{(j+1)} = 0, \quad j = 0, \dots, k-1, \quad (5.6)$$

are Equation (5.2) applied block-wise. They are algebraic, carry no derivative, and are equivalent to $\hat{f}^{(j+1)} = \mathbb{G} \hat{f}^{(j)}$ on the kernel. The *equation row*

$$\sum_{j=0}^k \mathbb{M}_{a_j} \hat{f}^{(j)} - \mathbb{M}_r \mathbb{D}^{(0)}(x_s) \hat{f}^{(0)} = 0, \quad (5.7)$$

is the equation itself, with $\mathbb{M}_{a_j}$ the multiplication operator of Section 3.3 (it reduces to $a_j I$ for a constant coefficient, and lifts the row to $n+1$ qubits where the standard form lifts it) and the source entering as there. The intermediate derivatives are unknowns of w and are never substituted by $\mathbb{G}^j \hat{f}^{(0)}$, so the dense step Equation (5.3) never appears. The name is borrowed by analogy with descriptor systems, which likewise keep the intermediate quantities as unknowns beside the state; here the $k+1$ block rows are algebraic in the coefficients of w .

Definition 11 (Descriptor residual operator): Collecting the recurrence rows and the

equation row, the descriptor residual operator A_{desc} is the block matrix

$$A_{\text{desc}} = \begin{pmatrix} & -\tilde{\mathbb{D}} & & \tilde{\mathbb{B}} & \\ & & -\tilde{\mathbb{D}} & \tilde{\mathbb{B}} & \\ & & & \ddots & \ddots \\ & & & & -\tilde{\mathbb{D}} & \tilde{\mathbb{B}} \\ \mathbb{M}_{a_0} - \mathbb{M}_r \mathbb{D}^{(0)} & \mathbb{M}_{a_1} & \cdots & \cdots & \mathbb{M}_{a_k} \end{pmatrix} \quad (5.8)$$

with $(k+1)N$ columns, and $A_{\text{desc}} w = 0$ is the equation. Its first k block rows are the recurrences and its last is the equation row.

The padded block register. A gate-level state indexes its block by a register of $n_f = \lceil \log_2(k+1) \rceil$ qubits, which addresses 2^{n_f} blocks and not $k+1$. The state is therefore padded to $2^{n_f}N$ amplitudes by blocks $b > k$ whose rows are $c_{\text{pad}} \hat{f}^{(b)} = 0$ with $c_{\text{pad}} = 1$. No other row reads a padding block, so the padding decouples: the Hamiltonian is block diagonal, and the physical blocks' spectrum is unaffected. The constraint rows of Section 5.4 are added to these block rows.

The mass form ($k=1$). A first-order equation or system $z' = \mathbb{A}z$ with C components has in the standard form the rows $I_C \otimes \mathbb{G} - \mathbb{A} \otimes I_N$, one block per component. The equation already gives each derivative, so a derivative block would carry what is determined. The mass form uses the equation to eliminate it: multiply the rows on the left by $I_C \otimes \tilde{\mathbb{B}}$ and use $\tilde{\mathbb{B}}\mathbb{G} = \tilde{\mathbb{D}}$,

$$(I_C \otimes \tilde{\mathbb{B}})(I_C \otimes \mathbb{G} - \mathbb{A} \otimes I_N) = I_C \otimes \tilde{\mathbb{D}} - \mathbb{A} \otimes \tilde{\mathbb{B}}. \quad (5.9)$$

Both factors are banded and none holds $\mathbb{G}$. The name is the ODE literature's, $\tilde{\mathbb{B}}$ playing the mass matrix of $\tilde{\mathbb{B}}\dot{z} = \dots$. For $f' + 2f = 0$ it gives $\tilde{\mathbb{D}} + 2\tilde{\mathbb{B}}$, on N unknowns, where a chain would have $2N$ unknowns and eliminating the derivative through $\mathbb{G}$ would give the dense $\mathbb{G} + 2I$; for a scalar equation with a constant coefficient it is the ultraspherical method's own operator in the basis of second-kind polynomials. The mass form has the standard form's state (one block per component), initial state and readout, and differs from it by a banded invertible factor on the equation rows alone. It has no derivative block and so nothing to rescale (Section 5.6). First-order systems carry the Carleman lifts of Section 6 and the coupled fields of Section 5.4, so the mass form is the layout that the nonlinear problems use.

Both layouts rest on one fact, proved next.

Theorem 1 (Kernel equivalence). *Let R_{std} stack A_{std} with the constraint rows of Section 3.3. Let R_{desc} stack the descriptor's recurrence, equation and padding rows with the same constraints, each reading the block of the derivative it constrains (Section 5.4), and let R_{mass} be the mass form's stacked residual. Then*

$$\ker R_{\text{desc}} = \{ (p, \mathbb{G}p, \dots, \mathbb{G}^k p, 0) : p \in \ker R_{\text{std}} \}, \quad \ker R_{\text{mass}} = \ker R_{\text{std}}. \quad (5.10)$$

Proof. Write $Ep = (p; \mathbb{G}p; \dots; \mathbb{G}^k p; 0)$. Chain: the padding rows force the padding blocks of a kernel vector w to zero, and since $\tilde{\mathbb{B}}$ is invertible each recurrence row gives $w_{j+1} = \tilde{\mathbb{B}}^{-1} \tilde{\mathbb{D}} w_j = \mathbb{G} w_j$, so $w = Ew_0$. The recurrence rows vanish on Ep for any p , because $(\tilde{\mathbb{D}} - \tilde{\mathbb{B}}\mathbb{G})\mathbb{G}^j p = 0$; every other row of R_{desc} acts on Ep as the standard form's row of the same name acts on p , since the equation row substitutes $\mathbb{G}^j p$ for $\hat{f}^{(j)}$ and a constraint row reads the block that holds $\mathbb{G}^j p$. Hence $R_{\text{desc}}Ep$ vanishes exactly when $R_{\text{std}}p$ does, and E is injective because its first block is the identity. Mass form: R_{mass} is R_{std} with its equation rows multiplied on the left by the invertible $I_C \otimes \tilde{\mathbb{B}}$ (Equation (5.9)), and the constraint rows unchanged. $\square$

The theorem is about kernels, and it is exact. An exact kernel exists only where the solution lies in the span of the basis; elsewhere both residuals are full rank, and the ground states are the smallest right singular vectors of different matrices. They differ, the descriptor's measuring a state's residual against a norm that carries its derivatives (Section 5.5) and the mass form's weighting the rows by $\tilde{\mathbb{B}}$, and they converge to the same function as n grows, at rates that we measure and do not prove. On R2a the decoded field's error against the analytic solution is 8.9×10^{-15} in the descriptor and 2.8×10^{-11} in the standard form at $n = 7$ (Section 8.4). We therefore claim equal kernels and observe a common limit, not equal ground states.

Cost, in the same breath. The chain is not free. Its state has $2^{\lceil \log_2(k+1) \rceil} N$ amplitudes against the standard form's N : n_f more qubits, and up to twice the unpadded $(k+1)N$ when $k+1$ is not a power of two (a second-order equation carries f , f' and f'' in four blocks). The padding is inert but counted. The mass form grows nothing.

5.4 Constraints, Partial Differential Equations and Coupled Fields

Data constraints sit on the block that holds the quantity they constrain. A Dirichlet condition $f(x_z) = 0$ reads block 0 with the rank-one row $\mathbb{B}(x_z)$, as in the standard form. A Neumann condition $f'(x_m) = 0$ reads block 1 with the same row, $\mathbb{B}(x_m)\hat{f}^{(1)} = 0$, where the standard form must write $\mathbb{B}(x_m)\mathbb{G}|\psi\rangle$. The descriptor reads a derivative where it is stored. The regular datum of Section 3.3 reads block 0, so the field is decoded and scaled exactly as in the standard form.

On several axes the field is block 0, and each axis a owns a chain $\partial_a f, \dots, \partial_a^{k_a} f$ up to the highest order that the equation or a constraint takes along it. The recurrences use that axis's factors $\tilde{\mathbb{B}}_a, \tilde{\mathbb{D}}_a$, tensored with the identity on the other axes. The heat equation ($k_t = 1, k_x = 2$) carries four blocks, and Laplace's and the wave equation ($2+2$) five, padded to eight. A mixed derivative would need a block of its own and is not built, a limitation of the construction. A system of coupled first-order fields takes the mass form, one block per component. Initial and boundary *functions* enter both forms as datum slices, rows that pin a slice of the state to a known function. They are available to the standard form too, and neither form doubles the state of a linear partial differential equation (PDE), so the descriptor's advantage there is not one of state size.

5.5 The Gap of a Carried Derivative

Carrying the derivatives costs the descriptor something the standard form does not pay. Let E be the map of Theorem 1. On the states that satisfy the recurrences, the descriptor's Rayleigh quotient is

$$\frac{\|R_{\text{desc}}Ep\|^2}{\|Ep\|^2} = \frac{\|R_{\text{std}}p\|^2}{\sum_{m=0}^k \|\mathbb{G}^m p\|^2}, \quad (5.11)$$

because the recurrence rows vanish on Ep and the numerator is the standard form's. By the min-max principle, the descriptor's singular values are bounded above by the generalised ones of this pair. The norm in the denominator carries every derivative, so the gap is flat only where the equation controls each carried derivative: where no state p orthogonal to the solution has a small residual against a large derivative. The bound is an upper bound; flatness is not proved, and we measure it. The standard form measures the same states by $\|p\|$, against which their residuals are large, and its gap stays flat; its difficulty is the other side of the quotient, the norm of Equation (4.4).

Where it holds and where it fails. For a one-axis equation whose leading coefficient has no zero on $[-1, 1]$ the continuum a priori estimate, $\|f^{(k)}\| \leq (\|Lf\| + \|\text{lower order}\|)/\min |a_k|$, suggests that the quotient stays away from zero on the complement of the kernel, and we measure the gap flat: on R2a the descriptor's is 2.87 at $n = 7$ and 2.87 at $n = 8$. It fails in three ways.

1. *A leading coefficient that vanishes.* A polynomial of degree N can concentrate within $\mathcal{O}(N^{-2})$ of an endpoint, where the Chebyshev points cluster, and each derivative of it grows by a factor $\mathcal{O}(N^2)$ there. Against such a state the leading term $a_k p^{(k)}$ is small where the lower orders are small too, so the quotient is small against the top block. Unscaled, the gap falls as N^{-4} at a simple endpoint zero and as N^{-2} at an interior one; on R3a, whose leading coefficient is $1 - x^2$, it is 2.7×10^{-8} at $n = 8$.
2. *Two axes.* The equation constrains only a sum of top derivatives, $f_{xx} + f_{yy}$ or $f_t - \kappa f_{xx}$. A high-frequency state with $\mathbb{G}_x^2 p \approx -\mathbb{G}_y^2 p$, or with its weight in the block the equation weights least, has a small residual against large blocks, and only the datum rows penalise it. On R5b the unscaled gap is 2.5×10^{-7} at $n = 5$ per axis.
3. *The field block crowded out.* The ground state's weight on the field block is about $\|b\|/(\sum_m \|\mathbb{G}^m b\|^2)^{1/2}$, with b the solution's coefficients, which the derivatives' growth makes small. An initial state prepared on the field block then has an overlap bounded by that weight, and amplification costs the reciprocal of the overlap. On R2c, $f'' + 5f' + 400f = 0$, the unscaled gap is 7.5×10^{-6} and the field block's weight is 5.4×10^{-6} at $n = 7$.

The gap's collapse is not an artefact of the lift, the padding, the arithmetic precision, the constraints, the basis or the reference method: each was tested and ruled out, and a descriptor

Try to prove that the gap is flat for a regular one-axis equation (the lower bound on the quotient of Equation (5.11) uniformly in N). If proved, put the proof in Appendix

built in the ultraspherical basis collapses at the same rates where the leading coefficient vanishes. The tests and the mechanism in full are given in Appendix G.

5.6 Block Rescaling

The cure is to rescale the carried derivatives so that no block outweighs the equation's own scale. Carry the block of order m divided by ζ^m ,

$$\hat{f}^{(m)} \mapsto \zeta^{-m} \hat{f}^{(m)}, \quad m = 0, \dots, k, \quad (5.12)$$

that is, replace R_{desc} by $R_{\text{desc}}S$ with $S = \text{diag}(1, \zeta, \dots, \zeta^k) \otimes I_N$. Because S is diagonal and invertible, the kernel becomes $S^{-1} \ker R_{\text{desc}}$ and Theorem 1 survives. The field block has weight one, so the readout is unchanged, and the padding keeps its weight and still decouples. What changes is the Hamiltonian, $SR^\dagger RS$, a congruence and not a similarity: the spectrum moves, and by Equation (5.11) the denominator becomes $\sum_m \zeta^{-2m} \|\mathbb{G}^m p\|^2$, so ζ sets how much each carried derivative weighs against the residual.

In the circuit the rescaling is free. Scaling the columns of block m by ζ^m multiplies the coefficient of every term of the block encoding of Section 5.7 by ζ to the order of its column block. It adds no gate and no ancilla, and changes only the subnormalisation. A general diagonal equilibration, a different weight for each amplitude, would need a diagonal block encoding of its own at $\Theta(\dim w)$ gates; it is the reason we restrict to weights constant on a block, and it is excluded from every circuit in this thesis.

The a priori rule. ζ is computed from the equation's coefficients and the resolution, never from a solution or a spectrum. For a one-axis equation whose leading coefficient a_k has no zero on $[-1, 1]$, $\zeta = 2\rho$ with $\rho = (\bar{a}_j/\bar{a}_k)^{1/(k-j)}$, j the lowest order present and $\bar{a}$ the mean of $|a|$. For constant coefficients ρ is the geometric mean of the characteristic roots' magnitudes, the scale at which a solution $e^{\lambda x}$ has $\|f^{(m)}\| \approx |\lambda|^m \|f\|$, so dividing block m by ρ^m balances the blocks at the equation's own scale. The constants of the rule, the 2 here and the $\frac{1}{2}$ below, were fixed once by sweeping ζ over the regular and the singular benchmark problems, and are never refitted to a problem: the rule is then checked against the sweep's best ζ on every problem. The filter's cost stays within twice its least over a window of ζ **two to eight times wide**, so ζ need only be right to a factor of about two. Where a_k has a zero on the interval, $\zeta = N/2$, which makes the quotient of the concentrated states of Section 5.5 of order one. On several axes, $\zeta_a = N_a$ per axis, balanced across axes by the top-order coefficients (the wave equation $f_{tt} = c^2 f_{xx}$ takes (cN, N)).

What it restores. On every one-axis problem that the descriptor runs the gap no longer falls with N , and the field block regains its weight. It is the physical subspace's gap, the one the filter meets, since the padding alone would cap the whole spectrum's at one. On R2c the gap rises from 7.5×10^{-6} to **0.715** and the field block's weight from 5.4×10^{-6} to **0.774** at $n = 7$.

Appendix G: the rule's check against the grid's best ζ (within about $1.2 \times$

On R3a the gap is 2.27 at $n = 5$ and 3.62 at $n = 8$, against 2.7×10^{-8} unscaled: it rises slowly with n and does not fall.

What it does not remove. Three limits, each a finding.

1. *The norm grows with ζ .* Block m carries ζ^m , so $\|R\| \gtrsim \zeta^k \bar{a}_k$ and $\|H\| \gtrsim \zeta^{2k}$. With ζ constant this is nothing. With $\zeta \propto N$, which is what a simple zero requires, $\|H\| = \Theta(N^{2k})$ and, at a flat gap, the degree is about $\tilde{\Theta}(N^2)$ at second order, against the standard form's $\tilde{\Theta}(N^4)$ (Section 5.9).
2. *Two axes.* The rescaled gap still falls with N , far more slowly than unscaled: on R5b it is 0.253 at $n = 3$ per axis and 0.0942 at $n = 5$.
3. *A double zero is beyond any block rescaling.* If a_k vanishes to second order at an endpoint, the concentrated states have quotient $\mathcal{O}(N^{-8}\zeta^{2k})$, which is the unscaled gap's N^{-8} at $\zeta = 1$, so balancing needs $\zeta \propto N^{4/k}$, which is N^2 at $k = 2$, by when $\|H\| = \Theta(N^8)$, the standard form's. The paper's own Legendre panels R3c and R3d multiply the associated Legendre operator by $1 - x^2$ and so have exactly this. The descriptor runs the equivalent problem for $v = f/\sqrt{1 - x^2}$, whose leading coefficient has simple zeros, and we make no claim for the double zero.

This is the descriptor's version of the diagonal preconditioner of the ultraspherical method [15], which bounds the condition number of its banded system. Here the diagonal scales the state's columns and is constant on each block, which is why it costs nothing in the circuit.

5.7 The Block Encoding: One Flat Linear Combination

Every block of A_{desc} is a short sum of *bands*, each a diagonal followed by a truncated shift. The diagonal of $\tilde{\mathbb{D}}$ is $2c$ on column c , affine in the index, apart from column 1 where the normalisation of T_0 intervenes; $\tilde{\mathbb{B}}$ and a constant multiplication are constant apart from their ends; a variable coefficient's multiplication and the Galerkin truncation contribute a few more such exceptions. An affine diagonal is a sum of $n + 1$ Pauli- Z terms, and each exception is one reflection,

$$\text{diag}(c)_{c=0}^{N-1} = \sum_{j=0}^{n-1} 2^j \frac{I - Z_j}{2}, \quad R_m = I - 2|m\rangle\langle m|. \quad (5.13)$$

So A_{desc} is one flat LCU of $\mathcal{O}(n)$ terms,

$$A_{\text{desc}} = \sum_{t=1}^L c_t U_t, \quad U_t = (\text{field word}) \otimes \text{SHIFT}^{s_t} P_t, \quad \alpha_A = \sum_{t=1}^L |c_t|, \quad (5.14)$$

where the field word acts on the block register, SHIFT^{s_t} is the truncated shift by s_t , and P_t is a Z -string or a reflection. Where the diagonal is exactly affine, the sum is tight, $\alpha = \max_c |a + bc|$ when the end values share a sign, so the dominant term $\tilde{\mathbb{D}}$ is encoded at its own norm $2(N - 1)$

and its defects cost $\mathcal{O}(1)$. A shift of the N -dimensional register is not unitary; it is applied on a register widened by one overflow qubit, so that rows falling outside the interval land in the half that the final projection discards. That qubit is the shift's only ancilla.

Algorithm 1 compiles the terms into the circuit of Figure 3, and the walkthrough that follows it is by line.




Figure 3: The descriptor block encoding: one PREPARE/SELECT/PREPARE † over the whole residual.

Input: the blocks of A_{desc} as bands, the block weights ζ

Output: PREPARE and SELECT circuits; the subnormalisation α_A

Write every block as a sum of bands: a diagonal and a truncated shift;

Fit each diagonal as affine in the column index plus a few defects, Equation (5.13);

Merge terms with equal axis factors across blocks; expand the resulting field-register pattern in Pauli words;

Multiply the coefficient of each term by ζ^m for its column block m ;

$\alpha_A \leftarrow \sum_t |c_t|$; PREPARE loads $\text{sgn}(c_t)\sqrt{|c_t|/\alpha_A}$ on the joint selector index of term t ;

SELECT: controlled on each selector register, apply the field word, then on each axis the diagonal and the shift;

return PREPARE, SELECT, α_A ;

Algorithm 1: Compiling the rows of A_{desc} into PREPARE and SELECT.

The terms (lines 1–4). Lines 1 and 2 are the decomposition above. Line 3 gathers terms that share their axis factors across the rows and columns they join, so that the block register carries their pattern once, as a small matrix expanded in Pauli words; the recurrences' common $\tilde{\mathbb{D}}$ is paid for once and not k times. Line 4 is the whole of block rescaling.

PREPARE and SELECT (lines 5–6). The selectors are factored, one register for the field operation and, on each axis, one for the diagonal and one for the shift, each as wide as the number of distinct values its terms use. SELECT then applies each distinct operation once, controlled on its own register, and not once per term; this is what keeps the ancilla at $\mathcal{O}(\log n)$ (Section 5.8). Rank-one rows, the sources and the data constraints, are encoded

through the feature map with a rotation tree that loads the source’s coefficients: no dense state preparation appears anywhere. The constructions and the term-by-term accounting are given in Appendix G.4.

5.8 Cost of the Descriptor Encoding

Proposition 3 (Cost of the descriptor encoding). *For a descriptor residual of a fixed problem, the flat LCU of Equation (5.14) is an exact block encoding of the system rows at $\alpha_A = \sum_t |c_t|$, with $\mathcal{O}(\log n)$ ancilla and $\mathcal{O}(n^2)$ gates. For the second-order family R2,*

$$\lceil \log_2(n+4) \rceil + 8 \text{ ancilla}, \quad \mathcal{O}(n^2) \text{ gates.} \quad (5.15)$$

On a regular one-axis equation at fixed ζ , the subnormalisation of the system rows is $\alpha_A / \|A_{\text{desc}}\|_2 = 1 + \mathcal{O}(1/N)$.

The ancilla are the selector registers (the diagonal’s width is $\lceil \log_2(n+4) \rceil$, for the n gates Z_j , the identity and three reflections), the field and shift selectors, the overflow qubit, and the flags that project the field terms; every width other than the diagonal’s is fixed by the problem. The gate count follows because each distinct operation is applied once and the shift’s multi-controlled ladder is $\mathcal{O}(n^2)$. The proof is given in Appendix G.5. The count is verified for R2’s system rows at $n = 2$ to 12.

Optimality, hedged. An LCU whose SELECT indexes L distinct operations needs $\lceil \log_2 L \rceil$ selector qubits [12], and an affine diagonal has exactly $n+1$ non-zero Walsh coefficients, so an LCU that expands diagonals in Pauli strings needs $\Omega(\log n)$ of them. The descriptor’s ancilla count is therefore of the least order among exact LCU encodings of that kind. It is not claimed against every encoding.

What it measures. The ratio $\alpha / \|\cdot\|$ is the factor by which the encoding exceeds the least possible subnormalisation, and the filter’s degree is linear in it. The proposition is about the system rows; we measure the ratio for the whole stacked residual, constraints included. It tends to one on regular one-axis problems: $\alpha_R / \|R\|_2 = 1.04$ on R2a at $n = 8$ (a resource estimate). It stays near two where the leading coefficient vanishes, because $\zeta \tilde{\mathbb{D}}$, $\zeta^2 \tilde{\mathbb{B}}$ and $\zeta^2 a_2$ are then comparable and comparable terms add in ℓ^1 though each is tight: 2.12 on R3a at $n = 7$. On two axes it is 2.29 on R5b at $n = 4$ per axis. We state the limit one is with its scope: regular one-axis problems only. Beside it, the structured encoding of the standard form is at 7.74 on R2a and 3.7×10^4 on R3a, the latter being the lifted multiplier’s looseness of Section 4.4.

Interpretation. The standard form’s structured encoding needs $n-1$ selector ancilla for $\mathbb{G}$ alone, because U_{odd} is a sum of $N/2$ shifts (Section 4.2), so its ancilla grow as $\Omega(n)$; the descriptor’s grow as $\log n$. That is the difference between a dense operator and a banded one. The better scaling is real but not yet visible at the sizes that circuits can reach: for the whole

Appendix G
(app:pr
is still
a scaff-
fold:
write
the
ancilla
and
gate
count
and
the
1 +
 $\mathcal{O}(1/N)$
ratio
(the
dom-
inant
term
 $\zeta^j \tilde{\mathbb{D}}$
is en-
coded
at its

reflection on R2a the descriptor needs 12 ancilla at $n = 2$ and 13 at $n = 5$, against 10 and 13 for the standard form.

The gates per query go the other way, and we state it plainly. At the sizes reached the descriptor's query costs more than the structured standard form's, having more terms and a block register: for the reflection on R2a the transpiled two-qubit count is 25 914 against 8094 at $n = 2$ and 34 098 against 22 856 at $n = 5$. The ratio falls with n , but both counts are $\mathcal{O}(n^2)$, so the data do not show that it crosses one. The descriptor's advantage is therefore the number of queries, the filter's degree, and not the gates per query; its state is larger too, by the chain's block register. The next section is where the advantage lies.

5.9 Consequences for the Spectrum and the Filter

Every block of the descriptor is $\mathcal{O}(N)$ at fixed ζ , the recurrences' $\zeta^j \tilde{\mathbb{D}}$ the largest, so $\|R_{\text{desc}} S\| = \Theta(N)$ and $\|H_{\text{desc}}\| = \Theta(N^2)$ whatever the order k : the derivative is banded and its norm is that of $\tilde{\mathbb{D}}$ and no longer that of $\mathbb{G}^k$. With the encoding at the norm, $\alpha_H/\|H_{\text{desc}}\| \rightarrow 1$ on regular one-axis problems, from 59.9 for the standard form's on R2a at $n = 8$ to 1.09 for the descriptor's. Since Equation (2.6) gives the degree as $d = \Theta(\sqrt{\alpha_H/\Delta} \log 1/\varepsilon)$, the descriptor's rescaled gap, flat on a regular one-axis problem, gives

$$\|H_{\text{desc}}\| = \Theta(N^2), \quad \Delta = \Theta(1) \text{ (measured)} \implies d = \tilde{\Theta}(N) \quad (\text{a one-axis ODE whose leading coefficient is } \Theta(N^2)) \quad (5.16)$$

against $\tilde{\Theta}(N^{2k})$ for the standard form with the best exact encoding (Equation (4.4)). The scope is part of the claim. Where the leading coefficient vanishes, or on two axes, the rescaling's weights raise the norm and the degree is about $\tilde{\Theta}(N^2)$ or faster, and for a double zero the standard form's or worse (Table 4). Neither form is poly-logarithmic in N : $\tilde{\Theta}(N)$ is polynomial in N and exponential in n , and the claim is a lower exponent, not an exponential speed-up. The filter's degree is the minimax edge filter's, which prepares every comparison in this thesis, and the floor for polynomial filters in general is no stronger than what Lin and Tong's optimality supports [13].

Table 4: Where the descriptor's degree changes. Norms follow from the block weights, gaps are measured. The block weight ζ is the a priori rule of Section 5.6.

Case	ζ	$\ H\ $	Gap	Degree
Standard form, order k , constant coefficients	–	$\Theta(N^{4k})$	flat	$\tilde{\Theta}(N^{2k})$
Descriptor, one axis, a_k without a zero	fixed	$\Theta(N^2)$	flat	$\tilde{\Theta}(N)$
Descriptor, one axis, a_k with a simple zero, $k = 2$	$N/2$	$\Theta(N^4)$	flat	$\tilde{\Theta}(N^2)$
Descriptor, two axes	N_a per axis	$\Theta(N^4)$	falls	faster than $\tilde{\Theta}(N^2)$
Descriptor, a double zero	$\propto N^2$	$\Theta(N^8)$	–	the standard form's or worse

At $n = 8$ on R2a the estimated degree (a resource estimate) is 9.1×10^3 for the descriptor, and 1.2×10^{10} as built and 1.5×10^9 at the ideal-encoding bound for the standard form; the measured growth of the degree with N is in Section 8.3. We quote every comparison both ways,

as built and at $\alpha = \|R\|_2$, which contains neither form's encoding; where the order reverses, as it does on R3 at the ideal bound and on the heat equation at small n , the reversal is a finding and is reported as one (Sections 8.3 and 8.4).

Two further statements fix what the comparison does and does not say. First, the relative gap: the standard form's is $\Delta/\|H_{\text{std}}\| = 5.9 \times 10^{-15}$ at $n = 7$ on R2a, the descriptor's 2.8×10^{-6} . An eigensolver applied to H_{std} resolves its ground state only to about machine precision divided by that ratio, so at these sizes it cannot, while the singular value decomposition of R , whose relative scale is $\Theta(N^{-2k})$ and not $\Theta(N^{-4k})$, does, and every reference in this thesis is that decomposition. The consequence is the degree, not that the ground state is unresolvable. Second, one post-selected application of a block encoding of R succeeds on a Haar-random state of dimension D with mean probability $\|R\|_F^2/(D\alpha_R^2)$, which is $\Theta(1)$ in the descriptor against N^{-1} per axis in the structured standard form, whose dense $\mathbb{G}^k$ dominates the Frobenius norm. This compares the encodings and is not a cost of the preparation, which post-selects no single query.

5.10 Summary

The chapter answers the second question. The residual can be reformulated so that its Hamiltonian is banded, without changing its solution, and the reformulation has a price.

1. Keeping the banded factors $\tilde{\mathbb{B}}$ and $\tilde{\mathbb{D}}$ of $\mathbb{G}$ and carrying the derivatives as unknowns, as a chain for order two and above or as a mass form for first-order equations and systems, makes the residual banded with its kernel exactly the standard form's (Theorem 1).
2. Its block encoding is one flat LCU with $\lceil \log_2(n+4) \rceil + 8$ ancilla for R2's system rows, order-optimal among exact LCU encodings, and its subnormalisation tends to the norm on regular one-axis problems (Proposition 3). Its price is a larger state, by the chain's block register, and more gates per query at reachable n .
3. The Hamiltonian's norm is $\Theta(N^2)$, and for a one-axis equation whose leading coefficient does not vanish the filter degree is $\tilde{\Theta}(N)$ against the standard form's $\tilde{\Theta}(N^{2k})$ (Equation (5.16)). Neither is poly-logarithmic.
4. *Limit.* The carried derivatives erode the gap where the leading coefficient vanishes and on two axes. An a priori block rescaling restores it on one axis at no gate cost, at the price of a larger norm where the leading coefficient vanishes; on two axes the gap still falls, and a double zero is beyond any block rescaling.

Sections 8.3 and 8.4 test the cost and the gap on the benchmark problems. Carrying the intermediate derivatives instead of eliminating them turns the physics-informed residual from dense to banded with its kernel unchanged, so that the Hamiltonian's norm falls from $\Theta(N^{4k})$ to $\Theta(N^2)$; the price is a block register and a gap that the derivative blocks erode, which an a priori rescaling restores on one axis at no gate cost.

6 Polynomial Nonlinearity

6.1 Two Treatments

6.2 The Doubled Space: Exact Encoding, Degenerate Kernel

The doubled Hilbert space method can create degenerate kernels. QITE can't select the physical solution from amongst a set of degenerate ground states, so although it is possible to encode the solution to the original problem in the ground state of the Hamiltonian, it is not possible to extract it from the ground state of the descriptor Hamiltonian.

Proposition 4 (Imaginary time projects onto a degenerate kernel).

$$\lim_{\beta \rightarrow \infty} \frac{e^{-\beta H_{\otimes}} \varphi_0}{\|e^{-\beta H_{\otimes}} \varphi_0\|} = \frac{\Pi_{\ker H_{\otimes}} \varphi_0}{\|\Pi_{\ker H_{\otimes}} \varphi_0\|}, \quad \Psi = \psi \otimes \psi. \quad (6.1)$$

6.3 The Carleman Lift in Space–Time

$$\frac{du}{dt} = \sum_{p=1}^P F_p u^{\otimes p} \longrightarrow \frac{dz}{dt} = \mathbb{A}_K z, \quad z = (u, u^{\otimes 2}, \dots, u^{\otimes K}), \quad (6.2)$$

$$\rho = \frac{\|u_0\| \|F_2\|}{|\operatorname{Re} \lambda_1(F_1)|} < 1, \quad \text{truncation error} = O(\rho^K), \quad (6.3)$$

$$\begin{aligned} A_{\text{st}}^{\text{std}} &= \mathbb{G}_t \otimes I - \frac{T}{2} I \otimes \mathbb{A}_K, \\ A_{\text{st}}^{\text{desc}} &= \tilde{\mathbb{D}}_t \otimes I - \frac{T}{2} \tilde{\mathbb{B}}_t \otimes \mathbb{A}_K = (\tilde{\mathbb{B}}_t \otimes I) A_{\text{st}}^{\text{std}}, \end{aligned} \quad \xi = \frac{2t}{T} - 1. \quad (6.4)$$

6.4 Products and Couplings on a Circuit

Proposition 5 (The nodal fold is optimal).

$$\mathfrak{m}_1^{(p)} = 2^{p\ell/2} \mathbb{U}^\dagger \Delta_p \mathbb{U}^{\otimes p}, \quad \alpha(\mathfrak{m}_1^{(2)}) = 2 = \|\mathfrak{m}_1^{(2)}\|. \quad (6.5)$$

6.5 Fourier–Galerkin Fluids: Burgers and Navier–Stokes

$$u_t + u u_x = \nu u_{xx}, \quad F_1 = \operatorname{diag}(-\nu k^2), \quad (6.6)$$

$$\omega_t + J(\psi, \omega) = \nu \nabla^2 \omega, \quad \nabla^2 \psi = -\omega \text{ on } [0, 2\pi)^2, \quad \omega_0^{\text{TG}} = 2 \sin x \sin y. \quad (6.7)$$

6.6 Preparation Versus Projection

Table 5: Which nonlinear route prepares the solution, and which only projects onto a kernel that contains it.

Route	Kernel	What imaginary time returns	Precondition
Doubled space (Wu et al.; R6)	degenerate, dimension $\geq 2^{2n-1}$	the projection of φ_0 onto the kernel	none
Carleman lift (R8–R10), either form	one-dimensional near-kernel	the truncated solution	$\rho < 1$

remove
once
pre-
vious
sec-
tion
com-
plete

7 Preparing and Verifying the Ground State

7.1 Overview and Model Configuration

Pipeline schematic (TikZ). Problem specification (equation, constraints, variant, n) $\rightarrow$ residual assembly (standard form | descriptor form) $\rightarrow$ stacked-residual block encoding $U_R \rightarrow$ reflection $V \rightarrow$ walk $W \rightarrow$ GQSP imaginary-time filter $\rightarrow$ post-selection $\rightarrow$ readout. Alongside: the classical reference (SVD of R) and the three verification tiers (circuit simulation, exact emulation, resource estimate) attached to the stages they check.

Figure 4: Model configuration: from problem specification to decoded field.

7.2 Problems and Model Inputs

7.3 The Classical Reference

$$|\psi_{\text{ref}}\rangle = \arg \min_{\|\psi\|=1} \|R\psi\|_2, \quad \kappa(R^\dagger R) = \kappa(R)^2. \quad (7.1)$$

7.4 From Stacked Residual to Quantum Walk

$$R = \begin{pmatrix} R_1 \\ \vdots \\ R_L \end{pmatrix}, \quad (\langle j| \otimes \langle 0_a|) U_R (|0\rangle \otimes |0_a\rangle) = \frac{R_j}{\alpha_R}, \quad \alpha_R = \left(\sum_{j=1}^L \alpha_j^2 \right)^{1/2}, \quad H = R^\dagger R. \quad (7.2)$$

Table 6: The shared constraint layer: every constraint row, in either form, is one of four kinds.

Kind	Operator	Use
Zero (invariant)	$\mathbb{B}(x)$ on the chosen field or derivative block	$f(x_z) = 0, f'(x_m) = 0$
Collapse (regular)	$\mathbb{D}^{(0)}(x_s)$	constants, sources; lifting linear terms in the doubled space
Ratio (vector data)	$\mathbb{F} = y_a(I - a\rangle\langle a) - w\langle a , \quad w = \sum_{c \neq a} y_c c\rangle$	multi-component initial data (Carleman); coupled fields
Datum slice	$(\mathbb{S}_a(x^*) - \hat{g} \tau(x_s)^\top / v) \Pi_f$	inhomogeneous boundary or initial functions, without doubling

Table 7: The benchmark ladder. Each problem adds one capability to those below it; R2–R6 are the paper’s Figs 3–7.

	Problem	Reference	New capability
R1	$f' + 2f = 0, f(-1) = 1$	$e^{-2(x+1)}$	smallest end-to-end system
R2	Fig. 3: $\{1, 4, 4\}, \{1, -2, -3\}, \{1, 5, 400\}$	closed form	order 2; Dirichlet and Neumann invariants; stiffness
R3	Fig. 4: Legendre, $m = 0$ ($l = 4, 5$) and $m = 1$ ($l = 5, 6$)	P_l, P_l^1	variable coefficients and lifts; an exact kernel ($m = 0$)
R4	Fig. 5: four inhomogeneous equations	closed form	source injection via $\mathbb{D}^{(0)}$; Maclaurin sources
R5	Fig. 6: Laplace (2-D), heat and wave (1+1-D)	closed form	multi-axis tensors; invariants versus datum slices
R6	Fig. 7: $4f'' + 2f'^2 + f = 0; f'' - 2f^2 + x = 0$	$1 - x^2/8$; boundary-value solve	doubled space (kernel degeneracy)
R7	$u' = v, v' = -\omega^2 u, \omega = \pi$	$\sin, \cos$	two fields; ratio rows
R8	$y' = -y + y^2, y(0) = \frac{1}{2}$	$1/(1 + e^t)$	Carleman lift, both forms in time
R9	viscous Burgers, periodic, Fourier–Galerkin	Cole–Hopf; DNS	field-scale Carleman; convolution F_2 ; QFT fold
R10	2-D vorticity Navier–Stokes, periodic	Taylor–Green; DNS	the culmination

Proposition 6 (Stacked-residual reflection).

$$\begin{aligned} V &= U_R^\dagger (2\Pi_{a=0} - I) U_R, & V^\dagger &= V, & V^2 &= I, \\ \Pi_{\text{in}} V \Pi_{\text{in}} &= \frac{2H}{\alpha_R^2} - I \equiv H', & C(V) &= U_R^\dagger C(2\Pi_{a=0} - I) U_R. \end{aligned} \quad (7.3)$$

$$W = (2\Pi_{\text{in}} - I) V, \quad \Pi_{\text{in}} W^k \Pi_{\text{in}} = T_k(H'). \quad (7.4)$$

7.5 The Filter and Its Parameters

$$\Pi_{\text{in}} P(W) \Pi_{\text{in}} = \sum_{k=0}^d q_k T_k(H'), \quad P(z) = \sum_{k=0}^d q_k z^k, \quad \sup_{|z|=1} |P(z)| \leq 1. \quad (7.5)$$

$$f_\beta(x) = e^{-\beta(1+x)}, \quad f_\beta(H') = e^{-2\beta H/\alpha_R^2}, \quad q_k = (2 - \delta_{k0}) (-1)^k e^{-\beta} I_k(\beta). \quad (7.6)$$

Proposition 7 (Normalisation of the imaginary-time filter).

$$\sum_{k=0}^{\infty} |q_k| = 1. \quad (7.7)$$

$$\beta = \frac{\alpha_R^2}{2\Delta} \ln \frac{1}{\varepsilon_f}, \quad d \approx \sqrt{2\beta \ln(2/\varepsilon_t)}, \quad (7.8)$$

$$|\varphi_0\rangle \propto \bigotimes_{j=0}^{n-1} (|0\rangle + r^{2^j} |1\rangle), \quad \gamma = |\langle \psi_{\text{ref}} | \varphi_0 \rangle|. \quad (7.9)$$

$$p_G = \|f_\beta(H') \varphi_0\|^2 \rightarrow \gamma^2, \quad \text{gates} \approx \frac{\pi}{4\gamma} \left[2d g(U_R) + d g(\text{refl}) \right]. \quad (7.10)$$

Optional quantikz circuit: one GQSP ancilla; $R_0 \prod_{k=1}^d [C_0 W R_k]$ with $C_0 W$ the controlled walk built from U_R , $U_R^\dagger$ and the two reflections, only the reflections controlled (Proposition 6); post-selection on the ancilla and the block-encoding register. The R_k phases realise whichever filter is in force (minimax, primary; imaginary time, the paper's).

Figure 5: The GQSP filter circuit.

7.6 Verification and the Error Budget

Table 8: Operator-verification modes: what each compares, up to what size, and to what tolerance.

Mode	Comparison	Size	Tolerance
Exact	statevector per input column; $\alpha \Pi_{\text{out}} U \Pi_{\text{in}}$ against the independently assembled operator	$\lesssim 20$ qubits	max-abs $\leq 10^{-10} \ X\ $
Probe	16–32 Haar-random and structured in- puts against $X\psi/\alpha$ by sparse matrix- vector products	$\lesssim 28$ –30 qubits	relative $\leq 10^{-9}$
IR	compiled term IR against the residual assembly at every size; each term type exactly at small n	any n	$\leq 10^{-12}$
Approximation	published $\mathbb{G}$ only: entry error $\epsilon_{\mathbb{G}}(n)$ against the small-angle bound	exact sizes	reported, not gated

Table 9: Acceptance criteria for every state-preparation run.

Level	Metric	Acceptance
State	infidelity $1 - \langle \psi_{\text{ref}} \psi \rangle ^2$ in the form’s own space	$\leq \varepsilon = 10^{-6}$, and $\leq 10\times$ the predicted bound
Probability	measured p_G against $\ P(H')\varphi_0\ ^2$	relative difference $\leq 10^{-8}$
Field	decoded f_Q against the analytic solution (L^∞ , relative L^2)	$\leq \varepsilon_{\text{disc}} + \kappa_{\text{read}} \sqrt{\varepsilon_{\text{prep}}}$
T1 against T2	state difference	$\leq 10^{-10}$
Phases	realised against target polynomial	$\leq 10^{-12}$

$$\|f_Q - f\| \leq \underbrace{\|f_{\text{ref}} - f\|}_{\varepsilon_{\text{disc}}} + \kappa_{\text{read}} (\varepsilon_{\text{filter}} + \varepsilon_{\text{trunc}} + \varepsilon_{\text{phase}} + d \varepsilon_{\text{BE}}) + \varepsilon_{\text{sample}}. \quad (7.11)$$

7.7 Reading Out the Solution

$$f^*(x) = 4P_C(x) - P_G(x) - \frac{1}{2N} = \sqrt{\frac{p_G}{N}} \tau(x) \cdot \psi, \quad (7.12)$$

$$f_Q(x) = f(x_s) \frac{f^*(x)}{f^*(x_s)}, \quad \eta_e = \frac{f(x_s)^2 p_G}{N f^*(x_s)^2}. \quad (7.13)$$

7.8 Design of the Numerical Experiment

7.9 Software and Computation

8 Numerical Study

Opening

8.1 Reproducing Wu et al.

$$\eta_e = \|b\|^2 = \frac{N}{\pi} \int_{-1}^1 \frac{f(x)^2}{\sqrt{1-x^2}} dx. \quad (8.1)$$

Table 10: The paper’s printed scale factors against the exact solution’s $\|b\|^2$ (Equation (8.1)) and the paper-faithful ground state at the printed n . [†]: the ground state is not trusted (a degenerate kernel, or a gap ratio below 100). Classical computation.

Fig.	Problem	n	Printed η_e	Ideal $\ b\ ^2$	Ground state	Agreement
3b	R2a	3	1.29	1.29009	1.28731	✓
3c	R2b	3	32.47	32.47741	32.22636 [†]	✓ (one unit low)
3d	R2c	5	451.71	451.70962	451.52312	✓
4a	R3a	3	1.75	1.74512	1.74512	✓
4b	R3b	3	1.49	1.48779	1.48779	✓
4c	R3c	5	108.51	108.51142	109.70433 [†]	✓
4d	R3d	5	135.38	135.38360	137.46159 [†]	✓
5a	R4a	3	1.35	1.34649	1.34478	✓
5b	R4b	3	14.74	14.74797	14.72531	✓ (one unit low)
5c	R4c	4	0.82	0.82116	0.82521	✓
5d	R4d	4	77.51	77.51167	76.93737	✓
6b	R5a	3	5.21559	5.21653	5.24407	×
6d	R5b	4	603.863	598.86606	598.87966	×
6f	R5c	5	216.469	239.82712	239.82712	×
7b	R6a	2	3.75367	3.52344	3.52344 [†]	×
7c	R6b	2	0.07984	0.04383	0.04182 [†]	×

Four panels, the paper’s Figs 5a–5d (R4a–R4d): the analytic solution, the paper-faithful ground state (rounded constraint points, Maclaurin sources) and R4a–R4d solved with exact roots and Chebyshev sources, at the printed n ; inset or second row: L^∞ error against n for both variants. Source: notebook *10_paper_reproduction*. Classical computation.

Figure 6: Printed choices distort the paper’s solutions at the printed resolution.

Verification error against n for the published, structured and descriptor encodings: exact mode (filled markers), probe mode (open markers), IR mode (crosses), with each mode's tolerance as a dashed line; a second panel gives transpiled CNOT count per query against n with fitted exponents. Source: notebook 20_pillar1_operators.

Figure 7: Every encoding is verified to its mode's tolerance at every size tested.

Decoded field $f_Q(x)$ against the analytic solution for R1 and R2a (both forms): circuit simulation (T1, with shot-sampled readout) and exact emulation (T2) overlaid; second panel: the error budget of Equation (7.11) broken into $\varepsilon_{\text{disc}}$, $\varepsilon_{\text{filter}}$, $\varepsilon_{\text{trunc}}$, $\varepsilon_{\text{phase}}$ and $\varepsilon_{\text{sample}}$ against n . Source: notebook 30_pillar2_stateprep.

Figure 8: The prepared state reproduces the analytic solution to within the predicted error budget.

8.2 The Pipeline Is Faithful

8.3 Where the Cost Lies

Table 11: Where the cost lies: the paper’s regime under its published encoding and under the structured exact encoding, against the descriptor form, for Fig. 3b ($\{1, 4, 4\}$, $f(-1) = 0$). Tier: T2 where measured, T3 beyond.

Quantity	Standard, published	Standard, structured	Descriptor
$\ H\ $	$\Theta(N^8)$	$\Theta(N^8)$	$\Theta(N^2)$
Gap Δ	7.04 (flat)	7.04 (flat)	0.454 (flat)
$\alpha_{\mathbb{G}}/\ \mathbb{G}\ _2$	$\sim N^3/2$	≤ 2.12	no $\mathbb{G}$
$\alpha_H/\ H\ $	$\sim N^{12}$	$22 \rightarrow 60$	$5.3 \rightarrow 1.05$
Gates per query	$O(n^3)$, approximate	$O(n^2)$, exact	$O(n^2)$, exact
Ancilla	$2n + 3$ ($\mathbb{G}$ alone)	$n + O(\log n)$ per $\mathbb{G}$	$\lceil \log_2(n + 4) \rceil + 8$
Degree at $n = 8$ ($\varepsilon = 10^{-6}$)	8.2×10^{23}	1.2×10^{10}	5.6×10^3
Degree scaling	$\sim N^{10}$	$\tilde{\Theta}(N^4)$	$\tilde{\Theta}(N)$

Three log-log panels against $N = 2^n$ for Fig. 3b: (a) $\|H\|$ and α_H for the published, structured and descriptor encodings; (b) the gap Δ ; (c) the filter degree d , with fitted exponents ($\approx N^{10}$, $\tilde{\Theta}(N^4)$, $\tilde{\Theta}(N)$). Exact emulation (T2) as filled markers, resource estimates (T3) as open markers. Source: notebook `40_scaling_and_claims`.

Figure 9: The encoder wall and the Hamiltonian wall are separate gaps.

8.4 The Descriptor’s Gap and Its Rescaling

8.5 Nonlinear Problems

8.6 Case Study: Two-Dimensional Navier–Stokes

8.7 Summary of Findings

Case by case (R2a, R2c, R3a, R5b, ...), the descriptor's cost over the standard form's, as built and at the ideal bound $\alpha = \|R\|$, with the reversals marked (R3 at the ideal bound; the heat panel at small n). Optional. Source: notebook 40_scaling_and_claims.

Figure 10: Three representative problems: where the comparison, quoted as built, reverses at the ideal encoding bound.

Panels across the benchmark ladder: the gap Δ and γ^2 unscaled against rescaled; the a priori ζ against the best ζ found on a grid; where the rescaling stops (a vanishing leading coefficient, a second axis). Source: notebook 42_descriptor_gap.

Figure 11: Block rescaling restores an $O(1)$ gap on one axis, and stops at a double zero or on two axes.

Three columns (R2c, R3a, R5b or R5c): top row, the decoded solution against the analytic one; bottom row, error against n for the standard and descriptor forms (and, for R2c, with and without block rescaling). Exact emulation (T2). Source: notebook 40_scaling_and_claims.

Figure 12: Three representative problems, solved in both forms.

(a) R8: error against Carleman order K for several time resolutions n_t , both forms; (b) R6: projection fidelity against physical fidelity; (c) R9 Burgers: error per order against a/ν , diverging past $a/\nu \approx 5$ while the kernel stays unique and trusted. Source: notebook 50_navier_stokes.

Figure 13: Carleman truncation converges geometrically once time is resolved; the doubled space only projects; on Burgers the error tracks a/ν , not ρ .

(a) Taylor–Green and vortex-pair vorticity snapshots against DNS, with the tier marked (T2 small, T3 beyond); (b) the degree pairs, standard against descriptor time axis, narrowing as the modes grow; (c) the tensor-layout coupling’s gates per query against the matchings, to 316 modes. Source: notebook 50_navier_stokes.

Figure 14: Navier–Stokes is reached by exact emulation and a resource estimate, not by circuit simulation.

9 Discussion

9.1 Why the Descriptor Form Wins, and What It Costs

9.2 End-to-End Complexity

9.3 What the Reproduction Says About the Published Method

9.4 Where the Construction Applies

9.5 Feasibility on Fault-Tolerant Hardware

9.6 Limitations

remove
once
pre-
vious
sec-
tion
com-
plete

10 Conclusions and Future Work

10.1 Conclusions

10.2 Future Work

11 References

- [1] H.-Y. Wu, A. E. Paine, E. Philip, A. A. Gentile and O. Kyriienko, *Quantum algorithm for solving nonlinear differential equations based on physics-informed effective Hamiltonians*, 2025. arXiv: 2504.13174 [quant-ph].
- [2] O. Kyriienko, A. E. Paine and V. E. Elfving, ‘Solving nonlinear differential equations with differentiable quantum circuits,’ *Physical Review A*, vol. 103, p. 052 416, 2021. arXiv: 2011.10395 [quant-ph].
- [3] A. E. Paine, V. E. Elfving and O. Kyriienko, *Physics-Informed Quantum Machine Learning: Solving nonlinear differential equations in latent spaces without costly grid evaluations*, 2023. arXiv: 2308.01827 [quant-ph].
- [4] L. N. Trefethen, *Spectral Methods in MATLAB*. SIAM, 2000.
- [5] A. M. Childs and N. Wiebe, ‘Hamiltonian simulation using linear combinations of unitary operations,’ *Quantum Information and Computation*, vol. 12, pp. 901–924, 2012. arXiv: 1202.5822 [quant-ph].
- [6] A. Gilyén, Y. Su, G. H. Low and N. Wiebe, ‘Quantum singular value transformation and beyond: exponential improvements for quantum matrix arithmetics,’ in *Proceedings of the 51st Annual ACM SIGACT Symposium on Theory of Computing (STOC 2019)*, 2019, pp. 193–204. arXiv: 1806.01838 [quant-ph].
- [7] O. Kyriienko, ‘Quantum inverse iteration algorithm for programmable quantum simulators,’ *npj Quantum Information*, vol. 6, p. 7, 2020. arXiv: 1901.09988 [quant-ph].
- [8] S. Scali, C. Umeano and O. Kyriienko, *Quantum topological data analysis via the estimation of the density of states*, 2023. arXiv: 2312.07115 [quant-ph].
- [9] A. E. Paine, *Effective-Hamiltonian Quantum Solvers for Differential Equations: Alternative Constructions and Function Encodings*, 2026. arXiv: 2609.26330 [quant-ph].
- [10] S. A. Cuccaro, T. G. Draper, S. A. Kutin and D. P. Moulton, *A new quantum ripple-carry addition circuit*, 2004. arXiv: quant-ph/0410184 [quant-ph].
- [11] A. Barenco et al., ‘Elementary gates for quantum computation,’ *Physical Review A*, vol. 52, p. 3457, 1995.
- [12] S. Chakraborty, R. Hazra, X. Li, C. Shao, Q. Wang and Z. Zhang, *Quantum singular value transformation without block encodings: near-optimal complexity with minimal ancilla*, 2025. arXiv: 2504.02385 [quant-ph].
- [13] L. Lin and Y. Tong, ‘Near-optimal ground state preparation,’ *Quantum*, vol. 4, p. 372, 2020. arXiv: 2002.12508 [quant-ph].
- [14] Z. Cai, R. Lazarov, T. A. Manteuffel and S. F. McCormick, ‘First-Order System Least Squares for Second-Order Partial Differential Equations: Part I,’ *SIAM Journal on Numerical Analysis*, vol. 31, no. 6, pp. 1785–1799, 1994. DOI: 10.1137/0731091

- [15] S. Olver and A. Townsend, ‘A Fast and Well-Conditioned Spectral Method,’ *SIAM Review*, vol. 55, no. 3, pp. 462–489, 2013. DOI: 10.1137/120865458 arXiv: 1202.1347.
- [16] A. W. Harrow, A. Hassidim and S. Lloyd, ‘Quantum Algorithm for Linear Systems of Equations,’ *Physical Review Letters*, vol. 103, p. 150 502, 2009. DOI: 10.1103/PhysRevLett.103.150502
- [17] D. W. Berry, A. M. Childs, A. Ostrander and G. Wang, ‘Quantum Algorithm for Linear Differential Equations with Exponentially Improved Dependence on Precision,’ *Communications in Mathematical Physics*, vol. 356, pp. 1057–1081, 2017. arXiv: 1701.03684 [quant-ph].
- [18] A. M. Childs and J.-P. Liu, ‘Quantum Spectral Methods for Differential Equations,’ *Communications in Mathematical Physics*, vol. 375, pp. 1427–1457, 2020. DOI: 10.1007/s00220-020-03699-z
- [19] P. C. S. Costa, D. An, Y. R. Sanders, Y. Su, R. Babbush and D. W. Berry, ‘Optimal Scaling Quantum Linear-Systems Solver via Discrete Adiabatic Theorem,’ *PRX Quantum*, vol. 3, p. 040 303, 2022. DOI: 10.1103/PRXQuantum.3.040303
- [20] D. An, J.-P. Liu and L. Lin, ‘Linear Combination of Hamiltonian Simulation for Nonunitary Dynamics with Optimal State Preparation Cost,’ *Physical Review Letters*, vol. 131, p. 150 603, 2023. DOI: 10.1103/PhysRevLett.131.150603
- [21] S. Jin, N. Liu and Y. Yu, ‘Quantum simulation of partial differential equations: Applications and detailed analysis,’ *Physical Review A*, vol. 108, p. 032 603, 2023. arXiv: 2212.14703 [quant-ph].
- [22] D. Jennings, M. Lostaglio, R. B. Lowrie, S. Pallister and A. T. Sornborger, *The cost of solving linear differential equations on a quantum computer: fast-forwarding to explicit resource counts*, 2023. arXiv: 2309.07881 [quant-ph].
- [23] J.-P. Liu, H. Ø. Kolden, H. K. Krovi, N. F. Loureiro, K. Trivisa and A. M. Childs, ‘Efficient quantum algorithm for dissipative nonlinear differential equations,’ *Proceedings of the National Academy of Sciences*, vol. 118, no. 35, e2026805118, 2021. arXiv: 2011.03185 [quant-ph].
- [24] H. Krovi, ‘Improved quantum algorithms for linear and nonlinear differential equations,’ *Quantum*, vol. 7, p. 913, 2023. arXiv: 2202.01054 [quant-ph].
- [25] M. Lubasch, J. Joo, P. Moinier, M. Kiffner and D. Jaksch, ‘Variational quantum algorithms for nonlinear problems,’ *Physical Review A*, vol. 101, p. 010 301, 2020. DOI: 10.1103/PhysRevA.101.010301
- [26] R. D. Somma, Y. Subasi and D. Orsucci, ‘Quantum algorithms for systems of linear equations inspired by adiabatic quantum computing,’ *Physical Review Letters*, vol. 122, p. 060 504, 2019. DOI: 10.1103/PhysRevLett.122.060504
- [27] D. An and L. Lin, ‘Quantum Linear System Solver Based on Time-Optimal Adiabatic Quantum Computing and Quantum Approximate Optimization Algorithm,’ *ACM Transactions on Quantum Computing*, vol. 3, no. 2, p. 5, 2022. DOI: 10.1145/3498331

- [28] L. Lin and Y. Tong, ‘Optimal Polynomial Based Quantum Eigenstate Filtering with Application to Solving Quantum Linear Systems,’ *Quantum*, vol. 4, p. 361, 2020. DOI: 10.22331/q-2020-11-11-361
- [29] L. Lin, *Lecture notes on quantum algorithms for scientific computation*, <https://arxiv.org/abs/2201.08141>, Jan. 2022.
- [30] C. A. Williams, S. Scali, A. A. Gentile, D. Berger and O. Kyriienko, *Addressing the Readout Problem in Quantum Differential Equation Algorithms with Quantum Scientific Machine Learning*, 2024. arXiv: 2411.14259 [quant-ph].
- [31] G. Brassard, P. Høyer, M. Mosca and A. Tapp, *Quantum Amplitude Amplification and Estimation*, 2000. arXiv: quant-ph/0005055.
- [32] D. W. Berry, A. M. Childs, R. Cleve, R. Kothari and R. D. Somma, ‘Exponential improvement in precision for simulating sparse Hamiltonians,’ in *Proceedings of the 46th Annual ACM Symposium on Theory of Computing (STOC 2014)*, 2014, pp. 283–292. arXiv: 1312.1414 [quant-ph].
- [33] Y. Suzuki, S. Uno, R. Raymond, T. Tanaka, T. Onodera and N. Yamamoto, ‘Amplitude estimation without phase estimation,’ *Quantum Information Processing*, vol. 19, p. 75, 2020. arXiv: 1904.10246 [quant-ph].
- [34] D. Grinko, J. Gacon, C. Zoufal and S. Woerner, ‘Iterative Quantum Amplitude Estimation,’ *npj Quantum Information*, vol. 7, p. 52, 2021. arXiv: 1912.05559 [quant-ph].
- [35] H. Nishi, T. Kosugi, Y. Huang, K. Hirose, Y. Okayama and Y.-i. Matsushita, *Quantum State Readout via Overlap-Based Feature Extraction*, 2025. arXiv: 2505.08613 [quant-ph].
- [36] D. Camps, R. Van Beeumen, L. Lin and C. Yang, ‘Explicit Quantum Circuits for Block Encodings of Certain Sparse Matrices,’ *SIAM Journal on Matrix Analysis and Applications*, vol. 45, no. 1, pp. 801–827, 2024. arXiv: 2203.10236 [quant-ph].
- [37] C. Sünderhauf, E. Campbell and D. Camps, ‘Block-encoding structured matrices for data input in quantum computing,’ *Quantum*, vol. 8, p. 1226, 2024. arXiv: 2302.10949 [quant-ph].
- [38] Z. Zhang and C. Shao, *Low-ancilla block encodings via Hamiltonian simulation*, 2026. arXiv: 2607.01843 [quant-ph].
- [39] D. Motlagh and N. Wiebe, ‘Generalized Quantum Signal Processing,’ *PRX Quantum*, vol. 5, p. 020368, 2024. arXiv: 2308.01501 [quant-ph].
- [40] A. Klappenecker and M. Rötteler, ‘Discrete cosine transforms on quantum computers,’ *Proc. SPIE 4740, Image and Signal Processing and Analysis (ISPA 2001)*, 2001. arXiv: quant-ph/0111038.
- [41] *QRTlib*, 2025. arXiv: 2510.16625 [quant-ph].
- [42] *Dyadic-Order Quantum Fractional Transforms*, 2026. arXiv: 2604.09295 [quant-ph].
- [43] D. Maslov, ‘Advantages of using relative-phase Toffoli gates with an application to multiple control Toffoli optimization,’ *Physical Review A*, vol. 93, p. 022311, 2016. arXiv: 1508.03273 [quant-ph].

- [44] M. A. Nielsen and I. L. Chuang, *Quantum Computation and Quantum Information*, 10th Anniversary. Cambridge University Press, 2010.
- [45] B. K. Berntson and C. Sünderhauf, ‘Complementary polynomials in quantum signal processing,’ *Communications in Mathematical Physics*, vol. 406, p. 161, 2025. arXiv: 2406.04246.
- [46] L. Laneve, *Generalized Quantum Signal Processing and Non-Linear Fourier Transform are equivalent*, 2025. arXiv: 2503.03026 [quant-ph].
- [47] H. Ni, R. Sarkar, L. Ying and L. Lin, *Inverse nonlinear fast Fourier transform on $SU(2)$ with applications to quantum signal processing*, 2025. arXiv: 2505.12615 [quant-ph].
- [48] J. M. Martyn, Z. M. Rossi, A. K. Tan and I. L. Chuang, ‘Grand Unification of Quantum Algorithms,’ *PRX Quantum*, vol. 2, p. 040 203, 2021. arXiv: 2105.02859 [quant-ph].
- [49] G. H. Low and I. L. Chuang, ‘Hamiltonian Simulation by Qubitization,’ *Quantum*, vol. 3, p. 163, 2019. arXiv: 1610.06546 [quant-ph].
- [50] L. Grover and T. Rudolph, *Creating superpositions that correspond to efficiently integrable probability distributions*, 2002. arXiv: quant-ph/0208112.
- [51] M. Möttönen, J. J. Vartiainen, V. Bergholm and M. M. Salomaa, *Transformation of quantum states using uniformly controlled rotations*, 2004. arXiv: quant-ph/0407010.
- [52] L. N. Trefethen, *Approximation Theory and Approximation Practice*. SIAM, 2013.
- [53] R. J. LeVeque, *Finite Difference Methods for Ordinary and Partial Differential Equations: Steady-State and Time-Dependent Problems*. SIAM, 2007.
- [54] J. P. Boyd, *Chebyshev and Fourier Spectral Methods*, 2nd. Dover Publications, 2001.
- [55] C. Canuto, M. Y. Hussaini, A. Quarteroni and T. A. Zang, *Spectral Methods: Fundamentals in Single Domains*. Springer, 2006.
- [56] R. G. Voigt, D. Gottlieb and M. Y. Hussaini, *Spectral Methods for Partial Differential Equations*. SIAM, 1984.
- [57] M. Abramowitz and I. A. Stegun, *Handbook of Mathematical Functions with Formulas, Graphs, and Mathematical Tables* (NBS Applied Mathematics Series 55). National Bureau of Standards, 1964.
- [58] P. B. Bochev and M. D. Gunzburger, *Least-Squares Finite Element Methods* (Applied Mathematical Sciences). Springer, 2009, vol. 166.
- [59] G. I. Taylor and A. E. Green, ‘Mechanism of the Production of Small Eddies from Large Ones,’ *Proceedings of the Royal Society of London. Series A, Mathematical and Physical Sciences*, vol. 158, no. 895, pp. 499–521, 1937. DOI: 10.1098/rspa.1937.0036
- [60] E. Hopf, ‘The partial differential equation $u_t + uu_x = \mu u_{xx}$,’ *Communications on Pure and Applied Mathematics*, vol. 3, no. 3, pp. 201–230, 1950. DOI: 10.1002/cpa.3160030302
- [61] J. D. Cole, ‘On a quasi-linear parabolic equation occurring in aerodynamics,’ *Quarterly of Applied Mathematics*, vol. 9, no. 3, pp. 225–236, 1951. DOI: 10.1090/qam/42889
- [62] Qiskit contributors, *Qiskit: An Open-source Framework for Quantum Computing*, 2024. DOI: 10.5281/zenodo.2573505

- [63] V. Bergholm et al., *PennyLane: Automatic differentiation of hybrid quantum-classical computations*, 2018. arXiv: 1811.04968.

Appendices

A Quantum Gate Reference	53
B Chebyshev and Ultraspherical Identities	53
B.1 The differentiation recurrence	53
B.2 The normalised basis	53
B.3 The Ultraspherical Connection	54
B.4 The $N = 4$ instance	55
C The Standard Form of Wu et al. in Full	56
C.1 The paper's printed matrices at $n = 2$	56
C.2 Lifts and folds	56
D The Published Circuits as Built	56
D.1 The circuits as built	57
D.2 The small-angle ladder and its error	57
D.3 Measured resources and error	58
E The Structured Exact Encoding of the Differentiation Matrix	59
E.1 Proof of Proposition 1	59
E.2 The arithmetic SELECT	59
E.3 Proof of Proposition 2	59
E.4 Subnormalisation across n	59
F The Shared Pipeline: Identities and Formulas	59
F.1 Proof of Proposition 6	59
F.2 Qubitisation: Chebyshev polynomials of the walk	59
F.3 The controlled walk	59
F.4 The imaginary-time filter	59
F.5 GQSP phases	60
F.6 The readout identity for probabilistic preparation	60

G The Descriptor Construction in Full	60
G.1 Block rows in full	60
G.2 Higher order, variable coefficients, and the padded block count	60
G.3 Proof of Theorem 1	61
G.4 The term IR and its compilation	61
G.5 Proof of Proposition 3	61
G.6 Block rescaling	61
H The Nonlinear Route in Full	61
H.1 The doubled space	61
H.2 Proof of Proposition 4	64
H.3 The Carleman lift	64
H.4 Ratio rows and the space-time forms	64
H.5 Product folds	64
H.6 Fourier–Galerkin operators	64
I Verification Protocol	64
I.1 Modes, sizes and tolerances	64
I.2 Structured probes	64
I.3 Per-atom residuals	64
I.4 The classical reference: eigh against SVD	64
I.5 The circuit-simulation memory and time model	64
I.6 Acceptance and the error budget	64
J Benchmark Catalogue	64
J.1 Linear problems	64
J.2 Nonlinear and coupled problems	64
J.3 Results by problem	64
K Reproduction Record	64
K.1 All panels	64
K.2 The Fig. 5 triage	64
K.3 Panels not reproduced, and printed errors	64
K.4 The paper’s claims	64
K.5 The paper’s filter and its parameters	64
K.6 The tau-truncation ablation	68
L Software and Reproducibility	68
L.1 Architecture	68
L.2 Records and provenance	68
L.3 Campaigns and computational resources	68
L.4 Tests	68

L.5 Regenerating the results 68

L.6 Environment 68

M Research Proposal 68

A Quantum Gate Reference

Table 12: The gates used in this thesis.

Gate	Circuit symbol	Matrix
Hadamard		
Pauli X, Y, Z		
Rotation $R_Y(\theta)$		
Phase $P(\phi)$		
CNOT		
Toffoli		
Multi-controlled X		
Zero-state reflection $2 0\rangle\langle 0 - I$		
Quantum Fourier transform		

B Chebyshev and Ultraspherical Identities

This appendix records the entries of the banded factors $\tilde{\mathbb{B}}$ and $\tilde{\mathbb{D}}$ introduced in Section 5.2, the exact factorisation $\mathbb{G} = \tilde{\mathbb{B}}^{-1}\tilde{\mathbb{D}}$, and the fully populated $N = 4$ instance of the first-order descriptor operator.

B.1 The differentiation recurrence

For $f = \sum_{\ell \geq 0} a_\ell T_\ell$ with derivative $f' = \sum_{\ell \geq 0} p_\ell T_\ell$ on $[-1, 1]$, the Chebyshev coefficients satisfy

$$c_{\ell-1} p_{\ell-1} - p_{\ell+1} = 2\ell a_\ell, \quad \ell \geq 1, \quad (\text{B.1})$$

with $c_0 = 2$ and $c_\ell = 1$ for $\ell \geq 1$; the $\ell = 1$ row therefore carries the factor $c_0 = 2$ on p_0 . Truncating to $N = 2^n$ coefficients and dropping the p_N term, Equation (B.1) is the matrix identity $\tilde{\mathbb{B}}_{\text{raw}} p = \tilde{\mathbb{D}}_{\text{raw}} a$ with $\tilde{\mathbb{D}}_{\text{raw}}$ a single superdiagonal, $(\tilde{\mathbb{D}}_{\text{raw}})_{\ell, \ell+1} = 2(\ell + 1)$, and $\tilde{\mathbb{B}}_{\text{raw}}$ two diagonals, $(\tilde{\mathbb{B}}_{\text{raw}})_{\ell, \ell} = c_\ell$, $(\tilde{\mathbb{B}}_{\text{raw}})_{\ell, \ell+2} = -1$.

B.2 The normalised basis

The construction of Section 5.2 works in the normalised basis $\varphi_0 = T_0/2^{n/2}$, $\varphi_\ell = T_\ell/2^{(n-1)/2}$ ($\ell \geq 1$), orthonormal under the discrete N -point Chebyshev inner product $\langle T_k, T_\ell \rangle_N = \frac{N}{2} c_\ell \delta_{k\ell}$ [4]: since $N = 2^n$, this gives $\|\varphi_0\|_N = \|\varphi_\ell\|_N = 1$ exactly, so the coefficient vector of a unit-norm function is a unit vector. Writing $S = \text{diag}(\sqrt{c_\ell})$, the normalised factors are $\tilde{\mathbb{B}} = S \tilde{\mathbb{B}}_{\text{raw}} S^{-1}$ and $\tilde{\mathbb{D}} = S \tilde{\mathbb{D}}_{\text{raw}} S^{-1}$, that is

$$\tilde{\mathbb{D}}_{\ell, \ell+1} = 2(\ell + 1)\sqrt{c_\ell/c_{\ell+1}}, \quad \tilde{\mathbb{B}}_{\ell, \ell} = c_\ell, \quad \tilde{\mathbb{B}}_{\ell, \ell+2} = -\sqrt{c_\ell/c_{\ell+2}}, \quad (\text{B.2})$$

so $\tilde{\mathbb{D}}$ has the single non-zero diagonal $(2\sqrt{2}, 4, 6, 8, \dots)$ and $\tilde{\mathbb{B}}$ has main diagonal $(2, 1, 1, \dots)$ and second superdiagonal $(-\sqrt{2}, -1, -1, \dots)$. Both hold $\mathcal{O}(N)$ non-zero entries at any N .

The dense Chebyshev derivative matrix is their exact quotient,

$$\mathbb{G} = \tilde{\mathbb{B}}^{-1}\tilde{\mathbb{D}}, \quad (\text{B.3})$$

an algebraic identity rather than a numerical approximation: $\tilde{\mathbb{B}}$ and $\tilde{\mathbb{D}}$ are read directly off the same recurrence that defines $\mathbb{G}$ (Equation (B.1)), so the equality holds exactly at every N , independent of whichever values a test suite happens to check. Appendix I reports this identity's residual against the reference implementation, alongside the rest of the composition-algebra checks.

B.3 The Ultraspherical Connection

The banded factorisation above is the first-derivative case of the ultraspherical spectral method of Olver and Townsend [15]. Stating the correspondence precisely is what fixes what this thesis adds to it.

Olver and Townsend represent a function's derivative not in the Chebyshev basis T_k but in the ultraspherical (Gegenbauer) basis $C_k^{(\lambda)}$, the family of polynomials orthogonal with respect to the weight $(1-x^2)^{\lambda-1/2}$; at $\lambda = 1$ these are the Chebyshev polynomials of the second kind, $C_k^{(1)} = U_k$. Differentiation of a Chebyshev series obeys

$$\frac{dT_k}{dx} = k C_{k-1}^{(1)}, \quad k \geq 1 \quad (\text{B.4})$$

[15, Eq. 2.5], so for $u = \sum_k u_k T_k$ the coefficients of u' in the $C^{(1)}$ basis are read off directly: $(\mathcal{D}_0 u)_k = (k+1)u_{k+1}$, a single superdiagonal, sparse at every truncation. Because u' 's coefficients now live in a different basis to u 's own, a second, purely algebraic operator is needed to convert a function's own Chebyshev coefficients into the $C^{(1)}$ basis without differentiating it: $\mathcal{S}_0$, with $(\mathcal{S}_0)_{k,k} = 1$ ($k = 0$), $\frac{1}{2}$ ($k \geq 1$), and $(\mathcal{S}_0)_{k,k+2} = -\frac{1}{2}$ [15, Eq. 2.8]. Comparing entries, $\tilde{\mathbb{D}}_{\text{raw}} = 2\mathcal{D}_0$ and $\tilde{\mathbb{B}}_{\text{raw}} = 2\mathcal{S}_0$ exactly: the recurrence $\tilde{\mathbb{B}}_{\text{raw}} p = \tilde{\mathbb{D}}_{\text{raw}} a$ of Equation (B.1) is literally $\mathcal{S}_0 p = \mathcal{D}_0 a$, doubled by the convention that folds the weight of T_0 into $c_0 = 2$ rather than into the polynomial itself. This is the discrete shadow the framing of Section 5.2 claims: the descriptor recurrence does not compute a new identity, it asserts an existing one – that p 's own Chebyshev coefficients, converted into the $C^{(1)}$ basis by $\mathcal{S}_0$, must equal the derivative computed directly by $\mathcal{D}_0$ – as one row of a linear system, rather than solving it for p via $\mathcal{S}_0^{-1}$.

The generalisation to order $k \geq 1$ in [15] raises the basis one level per derivative: a λ -th derivative needs coefficients in $C^{(\lambda)}$, obtained from the analogous pair $dC_k^{(\lambda)}/dx = 2\lambda C_{k-1}^{(\lambda+1)}$ ($k \geq 1$, giving $\mathcal{D}_\lambda$, still a single superdiagonal) and a conversion operator $\mathcal{S}_\lambda : C^{(\lambda)} \rightarrow C^{(\lambda+1)}$, still two diagonals [15, Eqs. 3.2–3.3]. To assemble one equation of order k , every term is raised to the common top basis $C^{(k)}$ by a chain $\mathcal{S}_{k-1} \cdots \mathcal{S}_{\lambda+1}$ applied after each term's own $\mathcal{M}_\lambda[a^\lambda]\mathcal{D}_\lambda$, giving one almost-banded operator $L = \mathcal{M}_k[a^k]\mathcal{D}_k + \sum_{\lambda=0}^{k-1} \mathcal{S}_{k-1} \cdots \mathcal{S}_\lambda \mathcal{M}_\lambda[a^\lambda]\mathcal{D}_\lambda$ that is solved once for the coefficients of $u^{(k)}$ [15, §3.2] – so a k -th order equation needs k distinct conversion

operators $\mathcal{S}_0, \dots, \mathcal{S}_{k-1}$ chained together and the lower-order unknowns eliminated.

The descriptor construction of Section 5.3 takes a different route to the same sparsity. Rather than raising the basis and eliminating, it keeps every derivative's coefficients in the *same* normalised Chebyshev basis, as its own block $\hat{f}^{(j)}$ of w , and reuses the single $\lambda = 0$ pair $(\tilde{\mathbb{B}}, \tilde{\mathbb{D}})$ between every consecutive pair of blocks (Equation (5.6)). No elimination is performed, and no operator beyond $\tilde{\mathbb{B}}, \tilde{\mathbb{D}}$ and the multiplication operators of Appendix G.2 is required at any order k – the price is the larger state w that carries every intermediate derivative explicitly instead of eliminating it (Appendix G.2). What this thesis adds to [15] is not a new sparsification of differentiation – that is their result – but applying the *same* $\lambda = 0$ identity block-wise to build an unreduced, differential-algebraic residual for the physics-informed Hamiltonian of [1], whose block structure is exactly what Section 5.7 block-encodes.

B.4 The $N = 4$ instance

At $n = 2$ ($N = 4$), Equation (B.2) gives

$$\tilde{\mathbb{B}} = \begin{pmatrix} 2 & 0 & -\sqrt{2} & 0 \\ 0 & 1 & 0 & -1 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}, \quad \tilde{\mathbb{D}} = \begin{pmatrix} 0 & 2\sqrt{2} & 0 & 0 \\ 0 & 0 & 4 & 0 \\ 0 & 0 & 0 & 6 \\ 0 & 0 & 0 & 0 \end{pmatrix}, \quad \mathbb{G} = \begin{pmatrix} 0 & \sqrt{2} & 0 & 3\sqrt{2} \\ 0 & 0 & 4 & 0 \\ 0 & 0 & 0 & 6 \\ 0 & 0 & 0 & 0 \end{pmatrix}. \quad (\text{B.5})$$

The top-right entry $3\sqrt{2}$ of $\mathbb{G}$ is the fill-in: coefficient 0 of the derivative depends on coefficient 3 of the input. Row ℓ of $\mathbb{G}$ is non-zero at columns $\ell + 1, \ell + 3, \dots$, so this is the extreme case – the lowest-order coefficient can depend on up to half the others, while the highest-order one (row 3 above) depends on none. The factorisation removes the fill-in entirely: no row of $\tilde{\mathbb{B}}$ or $\tilde{\mathbb{D}}$ touches more than two coefficients.

For the first-order equation $f' + 2f = 0$ of Section 5.3, $w = (\hat{f}^{(0)}; \hat{f}^{(1)})$ and $A_{\text{desc}} = \begin{pmatrix} -\tilde{\mathbb{D}} & \tilde{\mathbb{B}} \\ 2I_4 & I_4 \end{pmatrix}$ is the 8×8 matrix

$$A_{\text{desc}} = \left[\begin{array}{cccc|cccc} 0 & -2\sqrt{2} & 0 & 0 & 2 & 0 & -\sqrt{2} & 0 \\ 0 & 0 & -4 & 0 & 0 & 1 & 0 & -1 \\ 0 & 0 & 0 & -6 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ \hline 2 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 2 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 2 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 2 & 0 & 0 & 0 & 1 \end{array} \right]. \quad (\text{B.6})$$

The top four rows are the recurrence block $-\tilde{\mathbb{D}} \hat{f}^{(0)} + \tilde{\mathbb{B}} \hat{f}^{(1)} = 0$; the bottom four are the differential row $2\hat{f}^{(0)} + \hat{f}^{(1)} = 0$. Every non-zero entry lies on a diagonal or a near-diagonal, in place of the dense $\mathbb{G} + 2I$ of the eliminated form.

C The Standard Form of Wu et al. in Full

This appendix gives the $N = 4$ instance of the standard-form residual A_{std} that Section 3.3 builds for Wu et al.’s constant-coefficient case, and the entries of the lifting operator $\mathbb{M}_{x^p}$ that the same section uses for a variable coefficient.

The standard-form residual at $N = 4$. For Wu et al.’s stiff pilot equation [1], $f'' + 5f' + 400f = 0$, Equation (3.4) gives $A_{\text{std}} = \mathbb{G}^2 + 5\mathbb{G} + 400I$. At $N = 4$, using $\mathbb{G}$ (Appendix B),

$$A_{\text{std}} = \begin{pmatrix} 400 & 5\sqrt{2} & 4\sqrt{2} & 15\sqrt{2} \\ 0 & 400 & 20 & 24 \\ 0 & 0 & 400 & 30 \\ 0 & 0 & 0 & 400 \end{pmatrix}, \quad (\text{C.1})$$

verified directly against `pihm.chebyshev_derivative_matrix`. A_{std} is triangular here only because this equation has no odd-order term beyond f' at even N ; the off-diagonal entries sit three orders of magnitude below the diagonal, a large-coefficient conditioning problem stacked on top of $\mathbb{G}$ ’s own density.

The lifting operator $\mathbb{M}_{x^p}$. $\mathbb{M}_{x^p}$ is exact and *rectangular*: multiplying a degree- $(N-1)$ model by x^p can raise its degree by as much as p , so the result is represented one register wider, $2N \times N$, with no truncation. At $n = 1$ ($N = 2$, so $p \in \{1, 2\}$),

$$\mathbb{M}_{x^1} = \begin{pmatrix} 0 & \frac{1}{\sqrt{2}} \\ \frac{1}{\sqrt{2}} & 0 \\ 0 & \frac{1}{2} \\ 0 & 0 \end{pmatrix}, \quad \mathbb{M}_{x^2} = \begin{pmatrix} \frac{1}{\sqrt{2}} & 0 \\ 0 & \frac{3}{2\sqrt{2}} \\ \frac{1}{2} & 0 \\ 0 & \frac{1}{2\sqrt{2}} \end{pmatrix}, \quad (\text{C.2})$$

verified directly against `pihm.multiplication_matrix`. The general entry follows the same product-to-sum identity conjugated into the normalised bases of $|\tau(x)\rangle_n$ and $|\tau(x)\rangle_{n+1}$ that Section 3.3 states in closed form.

C.1 The paper’s printed matrices at $n = 2$

C.2 Lifts and folds

D The Published Circuits as Built

Section 3.4 describes, at the level a reader needs to follow the argument, the circuit Wu et al. use to load $\mathbb{G}$ ’s column-linear entries as a rotation amplitude [1]. This appendix gives the mechanism in full, together with a rigorous bound on its entrywise error that this thesis derives and verifies – the source paper states only that the error is “negligible” and decreases “exponentially” with n , without a closed form.

D.1 The circuits as built

The constraint rows (Figs 2 and 8). $\mathbb{B}(x)$ and $\mathbb{D}^{(0)}(x_s)$ are block-encoded by the inverse feature map $U_\tau^\dagger(x)$ followed by $\hat{S}_{n+1} = 2|0\rangle\langle 0| - I$ controlled inside a Hadamard sandwich on one ancilla, a two-term LCU on $n + 2$ qubits. The encoded block is $\mathbb{B}(x)/\sqrt{2^{n+1}}$, and $\mathbb{D}^{(0)}(x_s)/(\sqrt{2^{n+1}}/f(x_s))$ for the collapse row. The global phase of -1 in $\hat{S}_{n+1}$ is a relative phase once the circuit is controlled, and cannot be dropped.

The derivative circuit (Fig. 9). The registers are, in order, a flag qubit acted on by $R_Y(\omega_c)$ with $\omega_c = 2 \arccos \sqrt{1/3}$, an X ancilla, an offset register of $m = n$ qubits sandwiched between Hadamards, a value ancilla, and the n -qubit system. Two conditioned subroutines assign the entries on different bases. In $\hat{U}_R$ the offset $|i\rangle$ selects the $(2i+1)$ th superdiagonal and the system $|j\rangle$ the j th column; it consists of n singly-controlled $R_Y(\theta_k)$ gates on the value ancilla, controlled by the bits of $|j\rangle$, followed by n cyclic shifts $\hat{R}_n$. In $\hat{U}_C$ the offset selects the $(i+1)$ th row; it consists of n multi-controlled $R_Y(\varphi_k)$ gates, one shift $\hat{L}_n$ and $2n$ swaps. With $\mu = 2/\|\mathbb{G}\|_S$ the angles are $\theta_k = 2 \arcsin(-2^{k-1}\mu)$ and $\varphi_k = 2 \arcsin[-2^{k-1}\mu(1/\sqrt{2} - 1)]$. The circuit uses $m + n + 3 = 2n + 3$ qubits, its gate count is linear in n before the multi-controlled gates are decomposed and $O(n^3)$ after an $O(n^2)$ decomposition of each, and $\mathbb{G}$ is the top-left $2^n \times 2^n$ block of the unitary multiplied by $(2^{n-1} + 2^n)\|\mathbb{G}\|_S$. These are the paper's statements; their tests follow.

D.2 The small-angle ladder and its error

The mechanism. A ladder of n singly-controlled R_y rotations acts on one flag qubit, initialised to $|1\rangle$: bit m of the (register-arithmetic-shifted) column index controls $R_y(\theta_m)$,

$$\theta_m = 2 \arcsin(-2^m \mu), \quad m = 0, \dots, n-1, \quad (\text{D.1})$$

for a small real coefficient μ . Because rotations on the same qubit compose by exact angle addition ($R_y(a)R_y(b) = R_y(a+b)$, an $SO(2)$ identity), the net rotation for a column index p with bit set S is $R_y(\Theta(p))$, $\Theta(p) = \sum_{m \in S} \theta_m$, giving the flag qubit's $|0\rangle$ -amplitude

$$\text{raw}(p) = -\sin(\Theta(p)/2) = \sin\left(\sum_{m \in S} \arcsin(2^m \mu)\right). \quad (\text{D.2})$$

The row-0 correction (entries $\sqrt{2}p$ in place of $2p$) is a second, independently accumulated ladder on the same flag qubit, with per-bit coefficient $\mu(1/\sqrt{2} - 1)$, controlled on the same “this is row 0” flag the main construction already computes.

Why this is only approximate. $\text{raw}(p)$ in Equation (D.2) would equal the desired linear value μp exactly only if $\sin\left(\sum_m \arcsin(x_m)\right) = \sum_m x_m$, which does not hold in general – $\sin$ of a sum is not the sum of sines. The discrepancy is a genuine, quantifiable small-angle error.

Proposition 8 (Entrywise error of the ladder construction). *Let $J_3(p) := \sum_{m \in S} (2^m)^3$ for*

column index p with bit set S . Then

$$|\text{raw}(p) - \mu p| \leq \frac{\mu^3}{6} p^3 + O(\mu^5 N^5). \quad (\text{D.3})$$

Proof. Expand both non-linearities to leading order in the small parameter μ : $\arcsin(y) = y + y^3/6 + O(y^5)$ and $\sin(x) = x - x^3/6 + O(x^5)$. Summing the first expansion over the fired bits and substituting into the second,

$$\sum_{m \in S} \arcsin(2^m \mu) = \mu p + \frac{\mu^3}{6} J_3(p) + O(\mu^5 N^5), \quad (\text{D.4})$$

$$\begin{aligned} \sin\left(\sum_{m \in S} \arcsin(2^m \mu)\right) &= \sin(\mu p) + \cos(\mu p) \cdot \frac{\mu^3}{6} J_3(p) + O(\mu^5 N^5) \\ &= \mu p + \frac{\mu^3}{6} (J_3(p) - p^3) + O(\mu^5 N^5). \end{aligned} \quad (\text{D.5})$$

Since $p = \sum_{m \in S} 2^m$ and every $2^m \leq p$, the multinomial expansion of p^3 term-by-term dominates $J_3(p)$, with equality only when S is a single bit (p itself a power of two): $J_3(p) \leq p^3$. Taking the worst case $p = N - 1$ gives Equation (D.3). $\square$

This is a leading-order bound, valid once the largest single argument, $2^{n-1}\mu$, is small and well inside $\arcsin$'s domain – true from $n = 2$ onward at the default $\mu = 2/\|\mathbb{G}\|_S$ (Equation (3.11)), but not at $n = 1$, where $\mu \cdot 2^0 = \mu = 1$ exactly and the bound is not yet meaningful.

D.3 Measured resources and error

Table 13: The published encoding of $\mathbb{G}$ against the paper's claims. Classical exact-mode measurement.

n	N	$\ \mathbb{G}\ _S$	Qubits (claim $2n + 3$)	Gates	Bound on $\epsilon_{\mathbb{G}}$	Measured $\epsilon_{\mathbb{G}}$
2	4	6.15×10^1	—	—	—	—
3	8	1.18×10^3	—	—	—	—
4	16	2.04×10^4	—	—	—	—
5	32	3.38×10^5	—	—	—	—
6	64	5.50×10^6	—	—	—	—
7	128	8.88×10^7	—	—	—	—
8	256	1.43×10^9	—	—	—	—

E The Structured Exact Encoding of the Differentiation Matrix

E.1 Proof of Proposition 1

E.2 The arithmetic SELECT

E.3 Proof of Proposition 2

E.4 Subnormalisation across n

Table 14: Subnormalisation of the structured and published encodings of $\mathbb{G}$, relative to $\|\mathbb{G}\|_2$.

n	$\ \mathbb{G}\ _2$	Structured	Published
2	7.39	1.623	49.9
3	30.1	1.858	468
4	121	1.986	4.05×10^3
5	484	2.051	3.36×10^4
6	1935	2.084	2.73×10^5
7	7739	2.100	2.20×10^6
8	30 957	2.109	1.77×10^7
9	123 827	2.113	1.42×10^8
10	495 308	2.115	1.14×10^9

F The Shared Pipeline: Identities and Formulas

F.1 Proof of Proposition 6

F.2 Qubitisation: Chebyshev polynomials of the walk

F.3 The controlled walk

F.4 The imaginary-time filter

$$e^{-\beta x} = I_0(\beta) + 2 \sum_{k=1}^{\infty} (-1)^k I_k(\beta) T_k(x). \quad (\text{F.1})$$

F.5 GQSP phases

F.6 The readout identity for probabilistic preparation

G The Descriptor Construction in Full

G.1 Block rows in full

G.2 Higher order, variable coefficients, and the padded block count

For an order- k equation the state is $w = (\hat{f}^{(0)}; \dots; \hat{f}^{(k)})$ and A_{desc} stacks k recurrence block-rows $-\tilde{\mathbb{D}} \hat{f}^{(j)} + \tilde{\mathbb{B}} \hat{f}^{(j+1)} = 0$ above the single differential row $\sum_j \mathbb{M}_{a_j} \hat{f}^{(j)} = 0$: $k+1$ block-rows, each N coefficients wide, giving the classical $(k+1)N \times (k+1)N$ square matrix of Definition 11.

A variable coefficient $a_j(x)$ enters through $\mathbb{M}_{a_j}$, the square Galerkin truncation to the first N modes of the exact product (the multiplication of Section 3.3), built from its Chebyshev coefficients, or from its monomial coefficients converted exactly. $\mathbb{M}_{a_j}$ is banded, with bandwidth the degree of a_j , so it folds into the same shift-and-diagonal form as $\tilde{\mathbb{B}}$ and $\tilde{\mathbb{D}}$. The truncation is the only approximation in the linear construction; its error is confined to the top Chebyshev modes of a product that is smooth and resolved.

Padding to a power of two. The classical $(k+1)N \times (k+1)N$ matrix above is exact and complete as a linear system; it is not, however, the object any later stage of the pipeline actually builds. A gate-level realisation must index which block $\hat{f}^{(j)}$ a given operator acts on with a qubit register, and a register of m qubits addresses 2^m blocks, not $k+1$ of them for an arbitrary k . A_{desc} is therefore *embedded* in a $2^{\lceil \log_2(k+1) \rceil} N$ -dimensional space: the $k+1$ block-rows and block-columns above are kept exactly as derived, and $2^{\lceil \log_2(k+1) \rceil} - (k+1)$ further diagonal blocks are appended, each a row $c_{\text{pad}} I_N$ with $c_{\text{pad}} = 1$ pinning its own dummy field to zero, $\hat{f}_{\text{dummy}}^{(j)} = 0$; no other row reads a padding block, so the padding decouples. This padding is exact: the physical solution already satisfies the dummy rows trivially, since it has nothing in those blocks, so no error is introduced beyond the multiplication's truncation above. The qubit cost, however, is not uniform in k :

Order k	Blocks $k+1$	Field qubits $\lceil \log_2(k+1) \rceil$	Padded blocks $2^{\lceil \log_2(k+1) \rceil}$	Dummy blocks
1	2	1	2	0
2	3	2	4	1
3	4	2	4	0
4	5	3	8	3
7	8	3	8	0

so the state w the gate-level construction actually carries has $2^{\lceil \log_2(k+1) \rceil} N$ components, not $(k+1)N$: equal to the unpadded count exactly when $k+1$ is itself a power of two (orders

1, 3, 7, ...), and up to nearly double it otherwise – order 2, the most common case after order 1, already pays for a field register it only three-quarters uses. This is the qubit cost that must always accompany the gate-count comparison of Section 5.8.

G.3 Proof of Theorem 1

G.4 The term IR and its compilation

G.5 Proof of Proposition 3

G.6 Block rescaling

H The Nonlinear Route in Full

H.1 The doubled space

Section 3.5 states the published construction as quadratic only. A term of degree p in the unknown needs p copies, $|\Psi\rangle = |\psi\rangle^{\otimes p}$, of dimension N^p , with its p factors folded back into one Chebyshev series by $p - 1$ applications of the product rule (a cubic term by a tripled space and two folds, a quartic by a quadrupled space and three). The register therefore grows with the degree, a limitation noted in the concurrent work of Paine [9], and the argument of Proposition 4 should carry over, since the kernel of the corresponding Hamiltonian is again a subspace and the p -fold product states are not; this thesis does not use it beyond $p = 2$. The alternative it builds for higher degrees is the Carleman lift of Section 6.3.

Table 15: Accuracy of the classical reference, $1 - |\cos|$ against the exact coefficients, for the standard form by $\text{eigh}(R^\top R)$ and by SVD of R , and for the descriptor form.

n	Standard, eigh	Standard, SVD	Descriptor
2	1.7×10^{-1}	1.7×10^{-1}	1.6×10^{-2}
3	8.6×10^{-8}	8.6×10^{-8}	5.8×10^{-9}
4	0	$\leq 2 \times 10^{-16}$	$\leq 2 \times 10^{-16}$
5	3.7×10^{-15}	$\leq 2 \times 10^{-16}$	$\leq 2 \times 10^{-16}$
6	1.7×10^{-10}	$\leq 2 \times 10^{-16}$	$\leq 2 \times 10^{-16}$
7	1.2×10^{-4}	$\leq 2 \times 10^{-16}$	$\leq 2 \times 10^{-16}$
8	0.73	$\leq 2 \times 10^{-16}$	$\leq 2 \times 10^{-16}$

Table 16: The linear benchmark problems, R1–R5, with the paper’s printed invariant constraints.

	Equation	Invariant (printed)	Reference
R1	$f' + 2f = 0$	$-(f(-1) = 1 \text{ regular})$	$e^{-2(x+1)}$
R2a	$f'' + 4f' + 4f = 0$	$f(-1) = 0$	$\frac{1}{2}(1+x)e^{-2x}$
R2b	$f'' - 2f' - 3f = 0$	$f(0) = 0$	$\frac{1}{4}(e^{3x} - e^{-x})$
R2c	$f'' + 5f' + 400f = 0$	$f'(0) = 0$	$e^{-5x/2}(\cos \omega x + \frac{\sqrt{7}}{21} \sin \omega x),$ $\omega = \frac{15\sqrt{7}}{2}$
R3a	Legendre, $l = 4$	$f'(0) = 0$	P_4
R3b	Legendre, $l = 5$	$f(0) = 0$	P_5
R3c	associated Legendre, $l = 5, m = 1$	$f'(0) = 0$	P_5^1
R3d	associated Legendre, $l = 6, m = 1$	$f(0) = 0$	P_6^1
R4a	$(x-1)f'' - xf' + f = (x-1)^2$	$f'(-0.2) = 0$	closed form
R4b	$f'' + 4f' + 4f = e^{-2x}, \bar{p} = 3$	$f(-0.73) = 0$	closed form
R4c	$f'' - 5f' + 6f = xe^{2x}, \bar{p} = 5$	$f(0.91) = 0$	closed form
R4d	$f'' + 4f' + 20f = xe^{-2x} \sin 4x,$ $\bar{p} = 7$	$f'(-0.75) = 0$	closed form
R5a	Laplace, $f_{xx} + f_{yy} = 0$	$f(\pm 1, y) = f(x, -1) = 0$	$\cos \frac{\pi x}{2} \sinh \frac{\pi(y+1)}{2} / \sinh \pi$
R5b	heat, $f_t = \frac{1}{25}f_{xx}$	$f(t, \pm 1) = 0, f_x(t, \pm \frac{3}{4}) = 0$	$e^{-4\pi^2 t/25} \sin 2\pi x$
R5c	wave, $f_{tt} = 4f_{xx}$	$f(\frac{1}{8}, x) = f(\frac{7}{8}, x) = 0,$ $f(t, \pm 1) = 0$	$\cos 4\pi t \sin 2\pi x$

Table 17: The nonlinear and coupled benchmark problems, R6–R10.

	Equation	Data	Reference
R6a	$4f''' + 2f'^2 + f = 0$ (paper Eq. 29)	$f'(0) = 0, x_s = 0$	$1 - x^2/8$
R6b	$f''' - 2f^2 + x = 0$ (paper Eq. 34)	$f(0.02615) = 0$	boundary-value solve to 10^{-12}
R7	$u' = v, v' = -\omega^2 u, \omega = \pi$	$u(-1) = 0, v(-1) = \omega$	$u = \sin \omega(x+1)$
R8	$y' = -y + y^2, t \in [0, 3]$	$y(0) = \frac{1}{2} (\rho = 0.5)$	$1/(1+e^t)$
R9	viscous Burgers, periodic, $u_0 = a \sin x$	Fourier–Galerkin, $ k \leq K_c$	Cole–Hopf; DNS; exact lift
R10a	2-D vorticity Navier–Stokes, $\omega_0 = 2 \sin x \sin y$	$J \equiv 0$ (Taylor–Green)	exact viscous decay
R10b	2-D vorticity Navier–Stokes, vortex pair and resonant triads	active nonlinearity	DNS; exact lift

Table 18: State preparation by problem and form. Tier as marked.

	n	Infidelity (std)	Infidelity (desc)	Degree (std)	Degree (desc)	Tier
R1	—		—	—	—	
R2a	—		—	—	—	
R2b	—		—	—	—	
R2c	—		—	—	—	
R3	—		—	—	—	
R4	—		—	—	—	
R5	—		—	—	—	
R7	—		—	—	—	
R8	—		—	—	—	
R9	—		—	—	—	
R10	—		—	—	—	

H.2 Proof of Proposition 4

H.3 The Carleman lift

H.4 Ratio rows and the space-time forms

H.5 Product folds

H.6 Fourier–Galerkin operators

I Verification Protocol

I.1 Modes, sizes and tolerances

I.2 Structured probes

I.3 Per-atom residuals

I.4 The classical reference: eigh against SVD

I.5 The circuit-simulation memory and time model

I.6 Acceptance and the error budget

J Benchmark Catalogue

J.1 Linear problems

J.2 Nonlinear and coupled problems

J.3 Results by problem

K Reproduction Record

K.1 All panels

Table 19: The paper’s solutions at its printed resolution: L^∞ error of the paper-faithful and corrected variants, and the corrected variant’s best over n . Classical computation.

Fig.	Problem	Printed n	Faithful	Corrected	Corrected, best over n
3b	R2a	3	1.2×10^{-3}	1.2×10^{-3}	3.2×10^{-13} ($n = 5$)
3c	R2b	3	$1.8 \times 10^{-2\dagger}$	$1.8 \times 10^{-2\dagger}$	2.3×10^{-11} ($n = 5$)
3d	R2c	5	2.2×10^{-3}	2.2×10^{-3}	5.6×10^{-12} ($n = 6$)
4a	R3a	3	1.9×10^{-15}	1.9×10^{-15}	1.9×10^{-15} ($n = 3$)
4b	R3b	3	3.0×10^{-15}	3.0×10^{-15}	2.3×10^{-15} ($n = 4$)
4c	R3c	5	$8.7 \times 10^{-1\dagger}$	4.0×10^{-14}	4.9×10^{-15} ($n = 3$)
4d	R3d	5	$1.2^\dagger$	1.5×10^{-13}	8.0×10^{-15} ($n = 3$)
5a	R4a	3	3.1×10^{-3}	1.8×10^{-6}	3.8×10^{-14} ($n = 4$)
5b	R4b	3	1.9×10^{-2}	7.8×10^{-3}	8.5×10^{-12} ($n = 5$)
5c	R4c	4	7.2×10^{-3}	4.9×10^{-10}	1.6×10^{-13} ($n = 5$)
5d	R4d	4	6.1×10^{-2}	4.9×10^{-7}	5.1×10^{-12} ($n = 5$)
6b	R5a	3	3.9×10^{-3}	1.4×10^{-4}	2.1×10^{-13} ($n = 4$)
6d	R5b	4	1.9×10^{-4}	4.7×10^{-5}	4.6×10^{-12} ($n = 5$)
6f	R5c	5	8.5×10^{-8}	7.3×10^{-8}	6.1×10^{-11} ($n = 6$)
7b	R6a	2	$2.2 \times 10^{-16\dagger}$	$2.2 \times 10^{-16\dagger}$	–
7c	R6b	2	$6.8 \times 10^{-3\dagger}$	$6.8 \times 10^{-3\dagger}$	–

Table 20: Which printed choice causes the Fig. 5 error: the paper-faithful variant over $n = 4, \dots, 10$, each correction alone, and both together.

Fig.	Problem	Faithful	Exact root	Chebyshev source	Both, best
5a	R4a	3.1×10^{-3}	–	–	3.8×10^{-14} ($n = 4$)
5b	R4b	1.6×10^{-2}	4.0×10^{-2}	3.6×10^{-2}	8.5×10^{-12} ($n = 5$)
5c	R4c	7.2×10^{-3}	4.8×10^{-3}	1.0×10^{-2}	1.6×10^{-13} ($n = 5$)
5d	R4d	6.1×10^{-2}	8.1×10^{-2}	1.1×10^{-1}	5.1×10^{-12} ($n = 5$)

Table 21: The claims of Wu et al. tested in this thesis, and their verdicts.

	Claim	Verdict
A-01	$\mathbb{G}$, $\mathbb{M}$, $\mathfrak{n}_1$ and $\mathbb{B}$ as printed at $n = 2$	reproduced
A-02	the printed η_e of Figs 3–5 (11 panels)	within one unit of the last printed digit; 9/11 equal when rounded
A-03	the printed η_e of Fig. 6	not reproduced
A-04	the printed η_e of Fig. 7	not reproduced
A-05	the printed Laplace solution	a typographical error (cosh for cos)
A-06	the $\mathbb{B}$ circuit (Fig. 2): $n + 2$ qubits, prefactor $\sqrt{2^{n+1}}$	confirmed
A-07	the $\mathbb{G}$ circuit (Fig. 9): $2n + 3$ qubits, $O(n^3)$ gates, error decreasing exponentially	–
A-08	the $\mathbb{D}^{(0)}$ circuit (Fig. 8): prefactor $\sqrt{2^{n+1}}/f(x_s)$	confirmed
A-09	the linear H is positive definite with distinct eigenvalues	–
A-10	the doubled-space H has at least 2^{2n-1} zero eigenvalues	–
A-11	the time rule $t \geq \lambda_{\max}/(\lambda_2(n-1))$ suffices	–
A-12	the necessary terms scale as $O(\text{poly}(n))$ and are efficient	true for gates; not for the subnormalisation
A-13	the examples are solved accurately at the printed n	distorted where rounded points or Maclaurin sources enter

part and one for its odd part, of maximal degrees d^e and d^o , runs the two opposite-parity QSVT sequences and sums them by an LCU: a mixed-parity circuit. Each sequence uses the unitary dilation

$$U(\tilde{H}) = \begin{pmatrix} \tilde{H} & \sqrt{I - \tilde{H}\tilde{H}^\dagger} \\ \sqrt{I - \tilde{H}^\dagger\tilde{H}} & -\tilde{H}^\dagger \end{pmatrix}, \quad (\text{K.1})$$

built from PREPARE and SELECT subroutines over the unitary decomposition of $\tilde{H}$ with $m = \lceil \log_2 L \rceil$ ancilla for L unitaries. The time satisfies the approximate bound $t \geq \lambda_{\max}/(\lambda_2(n-1))$ (Equation (3.12)). For the paper’s two Legendre examples of its supplement, $t = 15$ and $t = 8$ recover the ground state with $(d^e, d^o) = (6, 7)$. Beyond the $O(1/\Delta)$ polylog($1/\varepsilon$) preparation time, the paper expects algorithmic thermalisation to reach $O(1/\sqrt{\Delta})$, and notes that GQSP, GQET and GQSVT could improve the circuit.

Table 22: Dropping the top k rows of the equation block (the Chebyshev tau method) gives an exact kernel with accuracy comparable to the paper’s least squares.

Problem	n	Least-squares L^∞	σ_0	Tau-truncated L^∞	σ_0
R2a	5	3.2×10^{-13}	3.0×10^{-11}	1.6×10^{-13}	0
R2a	6	2.2×10^{-12}	4.9×10^{-10}	3.8×10^{-12}	0
R2a	7	2.8×10^{-11}	7.8×10^{-9}	1.0×10^{-10}	0
R2b	5	2.3×10^{-11}	5.3×10^{-11}	3.7×10^{-12}	0
R2b	6	1.8×10^{-10}	8.5×10^{-10}	7.4×10^{-11}	0
R2b	7	9.8×10^{-10}	1.4×10^{-8}	8.3×10^{-10}	0
R2c	5	2.2×10^{-3}	1.6×10^{-2}	1.8×10^{-3}	0
R2c	6	5.6×10^{-12}	5.5×10^{-10}	1.8×10^{-11}	0
R2c	7	2.4×10^{-10}	1.1×10^{-8}	1.9×10^{-10}	0




Figure 15: The layers of the `pihm` package and the one direction in which they may import.

K.6 The tau-truncation ablation

L Software and Reproducibility

L.1 Architecture

L.2 Records and provenance

L.3 Campaigns and computational resources

Table 23: Resource classes for the Setonix campaigns.

Class	Request	Typical runs
S	work , 1 CPU, 4 GB, 30 min	construction, IR checks, small exact verification
M	work , 16 CPUs, 64 GB, 4 h	exact verification to 20 qubits; mid-size exact emulation
L	work , 64–128 CPUs, ~230 GB, 12–24 h	circuit-simulation ceilings; large emulation; probes to 30 qubits
XL	highmem , 1 TB	the largest emulation and probe runs

L.4 Tests

Table 24: The test tiers of the `pihm` package.

Tier	Content	Run
Unit and property	operator identities over random n and x ; primitives; arithmetic; IR; filters; phases	every commit
Golden	fixtures from the frozen legacy code and the paper’s printed $n = 2$ matrices	every commit
Reproduction	the paper’s targets, both variants	every commit
Integration	every problem, both forms, smallest n , all stages	every commit
Documentation	Sphinx build, doctests, spelling	every commit
Claims	one test per claim; writes the claims report	at gates
Slow	larger- n verification; circuit-simulation ceilings	at gates, on Setonix

L.5 Regenerating the results

L.6 Environment

M Research Proposal

Masters Research Proposal

Dhyan Patel [23340229]
The University of Western Australia

May 25, 2025

Title: Solving Nonlinear PDEs via Quantum Singular Value Transformation

Supervisor: Professor Jingbo Wang

1 Introduction

1.1 Aims

This thesis aims to develop and evaluate quantum algorithms for solving nonlinear partial differential equations (PDEs) using the Quantum Singular Value Transformation (QSVT) framework. Based on recent work that represents differential equations and their boundary conditions as effective Hamiltonians, this work extends the QSVT to the Generalised Quantum Singular Value Transformation (GQSVT). This extension enables the application of this method to complex-valued differential equations which arise naturally in many areas of physics and engineering. Additionally the project will explore how to incorporate data-informed modifications into these Hamiltonians, with the goal of modelling data-rich regimes.

1.2 Significance

Solving complex nonlinear partial differential equations (PDEs) on quantum hardware is a rapidly emerging prospect with broad applications ranging from fluid dynamics and turbulence modelling to quantum field theory. Established algorithms such as the Harrow-Hassidim-Lloyd (HHL) algorithm and variational methods typically rely on discretisation, which inhibits their scalability to stiff, nonlinear, multiscale or dense systems with fine grid sizes. In contrast, the effective Hamiltonian approach uses spectral methods to encode the solution into a Chebyshev polynomial basis, allowing it to surpass performance when fine-grid encoding is required. The quantum singular value transformation (QSVT) provides a flexible and modular approach to quantum linear algebra, allowing operations like matrix inversion under broader conditions than HHL [1]. However, the current method restricts effective Hamiltonians to real values, thereby limiting the class of operators and solutions to purely real outcomes. Moreover, the HHL paper only addresses linear partial differential equations or nonlinear ordinary differential equations

through a process called linearisation [3] and variational algorithms require classical optimisation [7].

This research is significant because it extends the effective Hamiltonian based framework to complex-valued nonlinear PDEs, which are common in frontier areas of various fields. In domains like geophysical inversion, deep turbulence and real-time weather prediction, incorporating data-informed terms into the Hamiltonian provides a promising path towards optimising quantum solutions to problems where complete physical models are unavailable or difficult to create. By enabling hybrid physics and data driven Hamiltonians, this work could expand the applicability of quantum PDE solvers to data-rich, real world problems.

2 Background

2.1 Problem Statement

Many physical systems are described by partial differential equations (PDEs), but solving them remains computationally costly. In the absence of closed-form solutions, many advanced classical algorithms implement numerical methods which are difficult to scale and lose accuracy in systems with a need for more fine grid or time-step discretisations. This is especially true of nonlinear PDEs which are the primary focus of this research.

Quantum algorithms provide a potentially scalable alternative. Quantum Singular Value Transformation (QSVT) allows quantum circuits to encode real-valued polynomials of definite parity. Recent theoretical developments have shown that it is possible to represent differential equations and their boundary conditions as effective Hamiltonians, which are structured matrices that can be processed using QSVT. However, due to QSVT's restriction to real-valued Hermitian Hamiltonians, the range of solvable PDEs is restricted. Systems with damping, gain/loss or complex boundary interactions can lead to complex-valued or non-Hermitian operators.

While novel frameworks such as the Generalised QSVT (GQSVT) have been proposed to handle complex-valued operators and remove parity constraints, they have not been systematically applied to nonlinear PDEs.

The existing effective Hamiltonian approach is largely physics-driven, assuming known governing models. Yet in many high-dimensional, data-rich domains such as geophysical imaging, real-time weather prediction, or turbulence modelling, the governing equations may be incomplete or unavailable. In these settings, data-informed Hamiltonian construction offers a promising path forward towards more flexible and accurate quantum solvers.

This research addresses the gap in current quantum PDE solving frameworks by developing and applying GQSVT based methods for nonlinear, complex-valued PDEs and exploring hybrid physics and data driven constructions for real world data-intensive applications.

2.2 Literature Review

2.2.1 Exploration of Problems

Burgers Equation (Viscous)

$$\frac{\partial u}{\partial t} + u \frac{\partial u}{\partial x} = \nu \frac{\partial^2 u}{\partial x^2} \quad (1)$$

This equation appears in turbulence theory and has a closed form solution in one real spatial dimension with $\nu > 0$. However, complex versions do exist and the equation remains unsolvable analytically outside these circumstances. The nonlinearity term is $u \frac{\partial u}{\partial x}$.

Nonlinear Schrödinger Equation

$$i \frac{\partial \psi}{\partial t} + \nabla^2 \psi + |\psi|^2 \psi = 0 \quad (2)$$

This equation describes the evolution of complex wave functions in nonlinear media. Its nonlinearity term is $|\psi|^2 \psi$.

Complex Ginzburg-Landau Equation

$$\frac{\partial A}{\partial t} = A + (1 + i\alpha) \nabla^2 A - (1 + i\beta) |A|^2 A \quad (3)$$

This equation models phenomena such as superconductivity and fluid convection patterns. The nonlinearity term is $|A|^2 A$.

The current physics-informed Hamiltonian method is able to address Burgers equation in its real case, but the Schrödinger and Ginzburg-Landau equations involve are complex-valued fields. This limitation motivates the development of a more robust framework capable of handling a broader class of nonlinear PDEs.

2.2.2 Existing Classical and Quantum Solving Methods

Classical methods for solving PDEs employ mesh based discretisation techniques such as finite difference and finite element methods or spectral methods [11]. These methods approximate the continuous domain onto a discrete grid, converting the PDE into a large system of algebraic equations that can be solved numerically. These methods are well-established and reliable, but for high-resolution or multi-dimensional systems the computational cost increases rapidly. For higher dimensional problems, direct matrix or classical variational methods may be employed [4].

The Harrow-Hassidim-Lloyd (HHL) algorithm was one of the first to demonstrate exponential speedup for solving systems of linear equations, which arise from discretised PDEs [3]. HHL encodes the linear system $A\vec{x} = \vec{b}$ into a quantum state and uses phase estimation to prepare a quantum state proportional to the solution vector $\vec{x}$. However, it is limited to sparse, well-conditioned Hermitian matrices. This makes it impractical for solving PDE problems, particularly nonlinear ones. Importantly, in the 2021 "Grand Unification" paper, it is shown that HHL can be viewed as a special case of QSVT, where

we're looking to find a target polynomial that inverts the singular values of the matrix A .

Variational quantum algorithms use a hybrid quantum-classical loop to optimise a parameterised quantum circuit that approximates the solution to a PDE [7]. This approach also discretises the domains into a grid. While more flexible than HHL, these methods require classical optimisation and suffer from challenges such as barren plateaus and scaling issues for larger systems.

2.2.3 Quantum Singular Value Transformation

The quantum singular value transformation (QSVT) is a powerful framework that generalises quantum linear algebra operations by enabling quantum circuits to apply polynomial transformations to the singular values of block encoded matrices [1]. QSVT enables the application of any parity-constrained polynomial $f(\Sigma)$ to the singular values Σ of a matrix A block encoded within a unitary.

$$U = \begin{bmatrix} A & \cdot \\ \cdot & \cdot \end{bmatrix} \text{ with } \|A\| \leq 1 \quad (4)$$

Assuming we have access to such a unitary U and the ability to apply a sequence of controlled-phase rotations by an angle ϕ , we can construct a unitary $V_{\vec{\phi}}$, dependent on the values of $\vec{\phi}$ we select for our phase rotations, that looks like the following:

$$V_{\vec{\phi}} = \begin{cases} \tilde{\Pi}_{\phi_1} U \left[\prod_{k=1}^{(d-1)/2} \Pi_{\phi_{2k}} U^\dagger \tilde{\Pi}_{\phi_{2k+1}} U \right] = \begin{bmatrix} \text{Poly}^{(\text{SV})}(A) & \cdot \\ \cdot & \cdot \end{bmatrix} & \text{for odd } \text{Poly}^{(\text{SV})}(A) \\ \left[\prod_{k=1}^{d/2} \Pi_{\phi_{2k-1}} U^\dagger \tilde{\Pi}_{\phi_{2k}} U \right] = \begin{bmatrix} \text{Poly}^{(\text{SV})}(A) & \cdot \\ \cdot & \cdot \end{bmatrix} & \text{for even } \text{Poly}^{(\text{SV})}(A) \end{cases} \quad (5)$$

where the sequences of phases $\vec{\phi} = (\phi_0, \phi_1, \dots, \phi_d)$ is chosen specifically to encode a desired polynomial $f(x)$ of degree d [1]. $\tilde{\Pi}_\phi$ and Π_ϕ represent a controlled phase rotation by ϕ where $\tilde{\Pi}$ and Π are the left and right unitary projectors respectively such that $A := \tilde{\Pi} U \Pi$. By carefully selecting which phase angles to rotate by, we can approximate any real, fixed parity continuous polynomial [6] acting on the singular values of A .

It can be shown that QSVT serves as a unifying framework for many quantum algorithms including quantum search, Hamiltonian simulation, factoring, quantum walks and more [8]. A key limitation of QSVT is that it only supports real-valued polynomials applied to matrices. This restricts its applicability to problems where both the matrix and the function being applied are real-valued [1]. For differential equations, this implies only PDEs with real-valued operators and real-valued solutions can be handled.

For physics-informed partial differential equation (PDE) solvers, QSVT is used in ground-state preparation via Hamiltonian simulation. In this case the target polynomial is $p(x) = e^{-tx}$ [11].




Figure 1: Quantum Circuit Diagram of the implementation of a Projector Controlled Phase Gate, Π_ϕ [2]

2.2.4 Physics-Informed Effective Hamiltonians

This algorithm encodes a differential equation in a latent Chebyshev basis and finds its solution by preparing the ground state of a specifically constructed Hamiltonian [11]. First, consider some differential equation with a differential operator $\mathcal{L}$, solution $u(x)$ and source function $f(x)$:

$$\mathcal{L}[u(x)] = f(x), \quad x \in [a, b]. \quad (6)$$

We can represent the solution as a state overlap in the Chebyshev basis with a scale factor s , namely

$$u(x) = s \langle T(x) | \psi \rangle \quad (7)$$

Here $|T(x)\rangle$ is known as a "feature map" state, a superposition of computational basis states with Chebyshev coefficients. $|\psi\rangle$ is a parameterised quantum state carrying the unknown amplitudes. Thus, we've effectively represented our solution in a latent Chebyshev basis [11].

To translate differential operators, we use the known differentiation matrices associated with the Chebyshev basis [5]. An example translation is shown below where $\mathbb{G}$ is the generator matrix corresponding to differentiation:

$$\frac{d^2 f(x)}{dx^2} + a f(x) = 0 \longrightarrow \mathbb{A} = \mathbb{G}^T + a \mathbb{I}. \quad (8)$$

Here, $\mathbb{A}$ is the matrix corresponding to the latent space representation of the DE. Crucially, differential operators act simply in the Chebyshev basis, as the following

$$\frac{d}{dx} f(x) = \frac{d}{dx} \langle T(x) | \psi \rangle = \langle T(x) | D^\top | \psi \rangle. \quad (9)$$

As a result, all terms in the DE can be rewritten as linear operators acting on $|\psi\rangle$.

To implement constraints, we use $\mathbb{B}(x)$, a data transformation operator dependent on x . For example, for specific constraints such as $f(x_0) = 0$ or $f'(x_1) = 0$, we can use $\mathbb{B}(x_0)$ or $\mathbb{B}(x_1)\mathbb{G}^\top$ [11]. From these, we can create our effective Hamiltonian:

$$\mathbb{H} = \mathbb{A}^\top \mathbb{A} + \mathbb{B}^\top \mathbb{B}. \quad (10)$$

We may add additional boundary constraints as required, which would effectively be more $\mathbb{B}$ matrices. Therefore, by construction, solving the original DE is equivalent to solving the eigen-problem:

$$\mathbb{H}_{\text{eff}}|\psi\rangle = 0, \quad (11)$$

and the desired solution corresponds to the ground state of the effective Hamiltonian. $\mathbb{H}$ is a real symmetric positive-definite matrix with distinct eigenvalues and mutually orthonormal eigenvectors [11].

We then block encode the effective Hamiltonian, and use a ground state preparation routine, QSVT, to amplify the amplitude of the ground state. Specifically, we use quantum imaginary time evolution, constructing the polynomial approximation of $e^{-t\mathbb{H}}$. This exponentially suppresses the excited eigen-components, leaving a higher probability of measuring the ground state. Importantly, the algorithm directly prepares the lowest energy solution without classical optimisation [11].

Finally, to measure the quantum state, we use a Chebyshev transformation and Hadamard test to measure the state we have prepared. This controlled overlap test effectively measures the overlap $\langle T(x)|\psi\rangle$. In practice this is done via interferometric measurement [11].

This same framework is extendable to multi-variable and nonlinear DEs. For a function of two variables $f(x, y)$, we use a tensor product Chebyshev basis with two registers, preparing a joint feature state $|T(x)\rangle \otimes |T(y)\rangle$. Take Laplace's equation for example:

$$\frac{d^2 f}{dx^2} + \frac{d^2 f}{dy^2} = 0 \longrightarrow \mathbb{A} = \mathbb{G}_x^{\top 2} \otimes \mathbb{G}_y^{\top 2}. \quad (12)$$

For nonlinear terms, we add extra registers. For example, a term like $f'(x)^2$ leads to a product of two copies of $|\psi\rangle$ in latent space. We create a new, higher dimensional latent Chebyshev space where $|\Psi\rangle \equiv |\psi\rangle \otimes |\psi\rangle$, so that the nonlinearity term becomes linear in the larger space. This naturally handles all nonlinearity. Inhomogeneous terms are handled via Maclaurin series expansion. For all these cases, we then construct the effective Hamiltonian as usual and solve [11]. This approach allows us to encode the solution over the entire domain, rather than using grid discretisation. Therefore, it is possible to evaluate the solution function at any arbitrary input values. However, as discussed above, QSVT can only handle real values. This restricts the class of differential equations we're able to solve.

Example Solutions:

Below are results from the application of this algorithm on example differential equations [11]. The paper did not attempt to solve nonlinear PDEs, but there are examples of both linear PDEs and nonlinear DEs. $f(x)$ is the known solution and $f_Q^*(x)$ is the quantum solution.

1. Heat Equation

$$\frac{\partial f(t, x)}{\partial t} - k \frac{\partial^2 f(t, x)}{\partial x^2} = 0, k = \frac{1}{25}, \mathcal{BC} = \{f(0, x) = \sin(2\pi x), f(t, \pm 1) = 0\} \quad (13)$$



Figure 2: Solution for the Heat Equation Example. [11]

2. Nonlinear Example 1

$$4 \frac{d^2 f(x)}{dx^2} + 2 \left(\frac{df(x)}{dx} \right)^2 + f(x) = 0, \mathcal{BC} = \{f(0) = 1, f'(0) = 0\} \quad (14)$$




Figure 3: Solution of Nonlinear Example 1 [11]

3. Nonlinear Example 2

$$\frac{d^2 f(x)}{dx^2} - 2f(x)^2 + x = 0, \mathcal{BC} = \{f(-1) = -0.1, f(1) = 0.1\} \quad (15)$$




Figure 4: Solution of Nonlinear Example 2 [11]

The model performs well for all examples shown. The errors are almost negligible and decrease as the number of qubits used (shown as n on the figures) increase. The experiments were conducted on classical simulations of quantum hardware and are therefore currently computationally costly, however based on these results it is projected this algorithm will become useful in the early fault-tolerant era [11].

2.2.5 Generalised Quantum Singular Value Transformation

The generalised QSVT (GQSVT) framework extends the standard QSVT algorithm by lifting the restrictions on coefficient phases and polynomial parity [9]. The standard QSVT restricts the target polynomial $p(x)$ to have real coefficients and definite parity. However, the generalised version allows complex polynomials and indefinite parity. These generalisations are inherited from developments in generalised quantum signal processing (GQSP) which is itself an extension of the quantum signal processing framework that QSVT was built on [10]. Eliminating the restrictions on real polynomials and Hamiltonians imposed by QSVT will allow us to solve complex-valued differential equations. GQSVT also promises a simplification in complexity for the crucial and time-consuming step of calculating phase factors. The task was of complexity $\mathcal{O}(d^2)$ for QSVT, GQSVT promises a complexity of $\mathcal{O}(d)$ [9]. This should speed up calculations even for problems which are currently possible via the existing method.




Figure 5: Example circuit for the generalised quantum singular value transformation for a Hermitian matrix A , block encoded as U . R is a rotation matrix analogous to Π for QSVT [9].

3 Research Plan

3.1 Method

This research will be primarily theoretical and computational. The goal is to design, simulate and evaluate novel quantum algorithms for solving nonlinear, complex-valued partial differential equations (PDEs). Development will be conducted using using *Mathematica* for symbolic manipulation and proof-of-concept derivations. Simulations will be created using *Python* based simulators such as *Qiskit* and *PennyLane* and results will be compared to classical solving algorithms. Efforts will be split into two phases. First, extending the current QSVT algorithm to GQSVT and second, integrating data-informed terms into effective Hamiltonians to enable solving PDEs in data-rich regimes. This research represents a new research direction within QUISA.

Proposed Plan:

1. Reproduce the Current Physics-Informed QSVT Method:
 - (a) Implement the standard QSVT framework in *Python*.
 - (b) Choose an appropriate problem.
 - (c) Implement a physics-informed Hamiltonian PDE solver in *Python*.
 - (d) Compare quantum solution against classical benchmarks
2. Extend to GQSVT:
 - (a) Create the GQSVT algorithm in *Python*.
 - (b) Choose an appropriate problem.
 - (c) Implement a new physics-informed Hamiltonian PDE solver in *Python*.
 - (d) Compare quantum solution against classical benchmarks
3. Experiment with more data-informed terms in the Hamiltonian

Timeline:

Semester 1 2025: Research Proposal.

Semester 2 2025: Recreation of QSVT Physics-informed method. Begin extension to GQSVT.

Semester 1 2026: Completion of a working GQSVT enhanced algorithm, including simulation and evaluation in comparison to classical algorithms.

Semester 2 2026: Experimentation with data-informed terms.

3.2 Preliminary Work

In semester 1 2025, work consisted of a thorough literature review, achieving an in-depth understanding of QSVT and the physics-informed Hamiltonian approach for PDEs. This was pivotal in developing a thesis that could both build on the newest developments in the field and be achievable within the time-frame of a Master's course. A working *Python* implementation of QSVT has been created and work is underway on implementing the current physics-informed Hamiltonian approach using QSVT.

References

- [1] András Gilyén, Yuan Su, Guang Hao Low, and Nathan Wiebe. Quantum singular value transformation and beyond: exponential improvements for quantum matrix arithmetics. In *Proceedings of the 51st Annual ACM SIGACT Symposium on Theory of Computing*, pages 193–204, New York, NY, USA, jun 23 2019. ACM. [Online; accessed 2025-05-18].

- [2] András Pál Gilyén. *Quantum Singular Value Transformation and Its Algorithmic Applications*. 2019. [Online; accessed 2025-05-25].
- [3] Aram W. Harrow, Avinatan Hassidim, and Seth Lloyd. Quantum algorithm for solving linear systems of equations. <https://arxiv.org/abs/0811.3171>, nov 19 2008.
- [4] Joshua Izaac and Jingbo Wang. *Computational Quantum Mechanics*. Springer, feb 15 2019. [Online; accessed 2025-05-25].
- [5] Chen Lan, Wei Li, and Huifang Geng. Pseudo-Hermitian Chebyshev differential matrix and non-Hermitian Liouville quantum mechanics. *Journal of Physics A: Mathematical and Theoretical*, 57(46):465302, nov 1 2024. [Online; accessed 2025-05-25].
- [6] Guang Hao Low and Isaac L. Chuang. Optimal Hamiltonian Simulation by Quantum Signal Processing. <https://arxiv.org/abs/1606.02685>, jun 8 2016.
- [7] Michael Lubasch, Jaewoo Joo, Pierre Moinier, Martin Kiffner, and Dieter Jaksch. Variational quantum algorithms for nonlinear problems. <https://arxiv.org/abs/1907.09032>, jul 21 2019.
- [8] John M. Martyn, Zane M. Rossi, Andrew K. Tan, and Isaac L. Chuang. Grand unification of quantum algorithms. *PRX Quantum*, 2(4), dec 3 2021. [Online; accessed 2025-05-16].
- [9] Danial Motlagh and Nathan Wiebe. Generalized Quantum Signal Processing. *PRX Quantum*, 5(2), jun 28 2024. [Online; accessed 2025-05-25].
- [10] Christoph Sünderhauf. Generalized Quantum Singular Value Transformation. <https://arxiv.org/abs/2312.00723>, dec 1 2023.
- [11] Hsin-Yu Wu, Annie E. Paine, Evan Philip, Antonio A. Gentile, and Oleksandr Kyriienko. Quantum algorithm for solving nonlinear differential equations based on physics-informed effective Hamiltonians. <https://arxiv.org/abs/2504.13174>, apr 17 2025.

Table 25: Deviations from the research proposal of May 2025.

Proposed	Outcome	Reason
Extending QSVT to the generalised construction for complex-valued, non- Hermitian operators Nonlinear PDEs: Burgers, nonlinear Schrödinger, com- plex Ginzburg–Landau Data-informed Hamiltonians (not proposed)		